


UCSD CSE 30

Computer Organization and Systems Programming

C Part 2

Paul Cao, Credit: Keith Muller



Variables in C

- Integer types
 - `char` [default: unspecified!]
 - `int` [default: signed]
- Floating Point
 - `float`, `double` [always signed]
- Optional Modifiers for each base type
 - `short` [int]
 - `long` [int, double]
 - `signed` [char, int]
 - `unsigned` [char, int]
 - `const`: read only
- char type
 - One byte in a byte addressable memory
 - Signed vs Unsigned implementation dependent
 - Be careful char is unsigned on arm and signed on other HW like intel

C Data Type	AArch-32 contiguous Bytes	AArch-64 contiguous Bytes
<i>%c</i> char (arm unsigned)	1	1
short int	<i>%hd</i> 2	2
unsigned short int	<i>%hu</i> 2	2
int	4 <i>%d</i>	4
unsigned int	4 <i>%u</i>	4
long int	<i>%ld</i> 4	8
long long int	<i>%lld</i> 8	8
<i>%f</i> float	4	4
<i>%lf</i> double	8	8
long double	<i>%llf</i> 8	16
pointer *	4	8

%p ↑ word size is the size of the address (pointer)

sizeof(): Variable Size (number of bytes) Operator

```
#include <stddef.h>
/* size_t type may vary by system but is always unsigned */
```

sizeof() operator returns a value of type **size_t**:

the number of bytes used to store a variable or variable type

```
size_t size = sizeof(variable_type);
```

or

```
size_t size = sizeof(variable_name); // preferred!
```

- The argument to sizeof() is often an expression:

```
size = sizeof(int * 10);
```

- reads as:

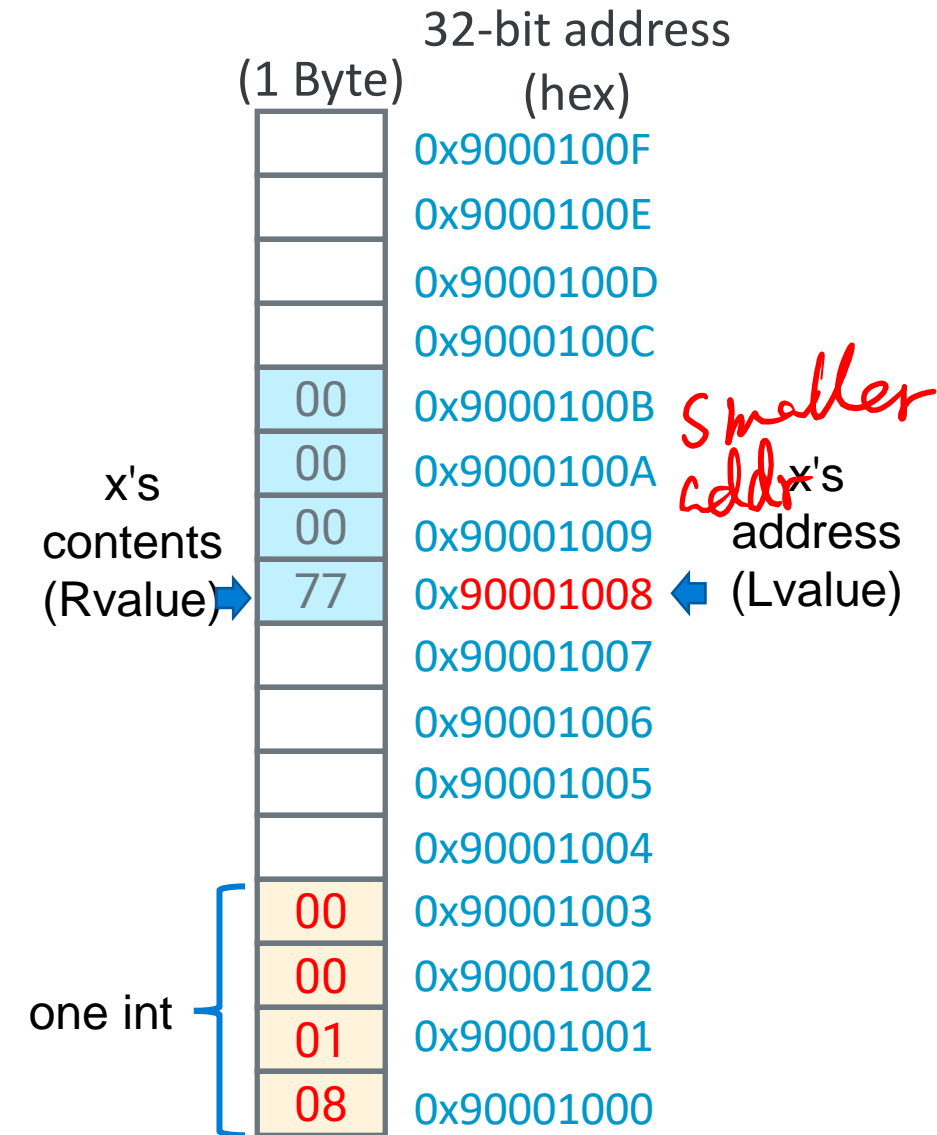
- **number of bytes** required to store **10 integers (an array of [10])**

Memory Addresses & Memory Content

```
Given int x, y; // 4 bytes on pi cluster
```

```
x = x; // Lvalue = Rvalue
```


- **Variable name** in a C statement evaluates to either:
 - **Lvalue:** when on the left side (Lside or Left value) of the = sign is the
 - address where it is stored in memory – a constant
 - Address assigned to a variable cannot be changed at runtime
 - **Rvalue:** when on the right side (Rside or Right value) of an = sign is the
 - contents or value stored in the variable (at its memory address)
 - requires a memory read to obtain



Memory Addresses & Memory Content

```
y = 42;
```

```
x = y;      // Lvalue = Rvalue
```



- **x** on left side (**Lside**) of the assignment operator = evaluates to:
 - The address of the memory assigned to the **x** – this is x's **Lvalue**
- **y** on right side (**Rside**) of the assignment operator = evaluates to:
 - **READ** the contents of the memory assigned to the variable **y** (type determines length – number of bytes) - this is y's **Rvalue**
- So `x = y;` is:

Read memory at y (**Rvalue**); write it to memory at x's address (**Lvalue**)

Introduction: Address Operator: &

call bit and

- Unary **address operator** (&) produces the **address** of where an **identifier** is in memory
- Requirement: **identifier must have a Lvalue**
 - Cannot be used with **constants** (e.g., 12) or **expressions** (e.g., x + y)
 - **&12** does not have an **Lvalue**, so **&12** is not a legal expression
- How can I get an **address for use on the Rside**? Three ways:
 - **&var** (any variable identifier or name)
 - **function_name** (name of a **function**, not func()); **&func_name** is equivalent
 - **array_name** (name of the **array** like array_name[5]); **&array_name** is equivalent

Introduction: Address Operator: &

- Unary **address operator** (&) produces the **address** of where an **identifier** is in memory

- **Example:** this might print:

value of g is: 42

address of g is: 0x71a0a0
(the address will vary)

```
int g = 42;
int
main(void)
{
    printf("value of g is: %d\n", g);
    printf("address of g is: %p\n", &g);
    return EXIT_SUCCESS;
}
```

Handwritten notes in the code block:
- Above the first printf: *int **
- Under the format specifier *%p*: *%p*
- Around the *&g*: *&g*

- **Tip:** printf() format specifier to display an address/pointer (in hex) is "%p"

Introduction: Pointer Variables - 1

- In C, there is a *variable type* for storing an address: a *pointer*
 - **Contents** of a pointer is an unsigned (0+, positive numbers) memory address
- When the **Rside** of a variable contains a **memory address**, (it **evaluates** to an **address**) the variable is called a **pointer variable**
- A **pointer** is defined by placing a *star* (or *asterisk*) (*) before the identifier (name)

```
type *name; // defines a pointer; name contains address of a variable of type
```

pointer


Introduction: Pointer Variables - 1

```
type *name; // defines a pointer; name contains address of a variable of type
```

- You also must specify the **type of variable** to which the pointer points

```
int i = 42;  
int *p; /* p contains the address of an integer */  
p = &i; /* p "points at" i (assign address of i to p) */
```

*int ** ptr;*



i ' 0x80

- Recommended:** be careful when defining multiple pointers on the same line:

`int *p1, p2;` is not the same as: `int *p1, *p2;`

Use instead:

```
int *p1;  
int *p2;
```

Introduction: Pointer Variables - 2

- Pointers are typed! Why?
 - The compiler needs the **size** (`sizeof()`) of the data **you are pointing at** (number of bytes to access)

- A pointer definition:

```
int *p = &i;  /* p points at i (assign address i to p) */
```

- Is the same as writing the following definition and assignment statements

```
int *p;      /* p is defined (not initialized) */  
p = &i;      /* p points at i (assign address i to p) */
```


- The ***** is part of the definition of **p** and is **not part of the variable name**
 - The name of the variable is **simply p**, not ***p**
- C mostly ignores whitespace, so these three definitions are equivalent

```
int  *p = &i;      /* Style A */  
int * p = &i;      /* Style B */  
int*  p = &i;      /* Style C */
```


Introduction: Pointer Variables - 3

- As with any variable, its value can be changed

```
p = &j;    /* p now points at j */
```



```
p = &i;    /* p now points at i */
```



Introduction: Pointer Variables - 4

- Pointer variables all use the **same amount of memory** no matter what they point at

```
int *iptr;  
char *cptr;  
  
printf("iptr(%u) cptr(%u)\n", sizeof(iptr), sizeof(cptr));
```


- Above prints on a 32-raspberry pi `iptr(4) cptr(4)`

Introduction: Indirection (or dereference) Operator: *

*
define
ptr
deref
mul-1

- The **indirection operator** (*) or the *dereference operator to a variable* is the **inverse** of the *address operator* (&)
- **address operator** (&) can be thought of as:

“get the address of this box”



- **indirection operator** (*) can be thought of as:


“follow the arrow to the next box and get its contents”

Introduction: Indirection (or dereference) Operator: *

*Contents of **p** is the **address** of **i** (p points at i)*

```
int i = 42;  
int j = i;  
int *p;  
p = &i;  
  
printf("*p is %d\n", *p);
```


```
% ./a.out  
*p is 42
```



Introduction: Indirection Operator Rside

- Performs the following steps when the ***** is on the Rside:
 1. **read** the **contents** of the **variable** to get an **address**
 2. **read** and return the contents at that address
 - (requires **two reads of memory on the Rside**)

```
z = *x; // copy the contents of memory pointed at by x to z
```




Introduction: Indirection Operator Lside

Performs the following steps when the ***** is on the Lside:

1. **read** the **contents** of the **variable** to get **an address**
2. **write** the evaluation of the Rside expression to that address
 - (requires **one read of memory and one write of memory on the Lside**)


```
*z = x; // copy the value of x to the memory pointed at by z
```




Each use of a * operator results in one additional read -1

Each * when used as a dereference operator in a **statement** (Lside and Rside) generates an additional read

```
int x = 2, y = 1;  
x = y; // one read
```




```
int z = 2, y = 1;  
int *x;  
x = &y;  
z = *x; // two reads
```




Each use of a * operator results in one additional read -2

- Each * when used as a dereference operator in a **statement** (Lside and Rside) generates an additional read

```
int z = 2, y = 1;  
int *x;  
x = &y;  
*x = z;
```



```
int z = 2, y = 1;  
int *x;  
int *w;  
x = &y;  
w = &z;  
*x = *w;
```



Recap: Lside, Rside, Lvalue, Rvalue

```
int x = 2, y = 1;  
x = y;
```

```
int z = 2, y = 1;  
int *x;  
int *w;  
x = &y;  
w = &z;  
*x = *w;
```

```
*x on Lside is 0x108  
w on Rside is 0x100  
*w on Rside is 2
```

Constant
Var Name

y

x

Constant
Var Name

x

y

z

w

Lvalue
address

0x108

0x104

Lvalue
address

0x10c

0x108

0x104

0x100

Rvalue
Contents

0x1

0x1

Rvalue
Contents

0x108

0x2

0x2

0x104

read

write

read (address)

write

read


read (address)

Pointer Practice

```
int *ptr;
```

Declares a variable, `ptr`, which is a pointer to (it contains the address of) an `int` in memory

`ptr`



```
int x = 5;
```

```
int y = 2;
```

Declares two variables, `x` and `y`, that contain ints, and *initializes* them to 5 and 2, respectively

`x`



write

`y`




write

```
ptr = &x;
```


Sets `ptr` to contain the address of `x` ("`ptr` points to `x`")

`ptr`

write



`x`



5

`y`



2

```
y = 1 + *ptr;
```

"Dereference `ptr`"


Sets `y` to "1 plus the value stored at the address held by `ptr`. Because `ptr` points to `x`, this is equivalent to `y = 1 + x;`"

`ptr`

read




`x`



5

read

`y`



6

write

```
x = *(&y);
```

Sets `x = y`; The `*` and `&` cancel each other. get the address of `y` and then get the contents pointed by that address

`ptr`



`x`



6

write

`y`



6

read

x

The NULL Constant and Pointers

- **NULL is a constant** that **evaluates to zero (0)**
- You **assign a pointer variable** to contain **NULL** to **indicate that the pointer does not point at anything**
- A **pointer variable** with a **value of NULL** is called a “**NULL pointer**” (invalid address!)
- Memory location 0 (address is 0) is not a valid memory address in any C program
- Dereferencing NULL at runtime will cause a program fault (segmentation fault)!

```
p = NULL;  
i = *p;          /* segmentation fault! */  
*(int *)900000 = 25; /* cast 900000 to a pointer */  
                    /* if writeable address space, it works */  
                    /* that memory location just changed */
```

Using the NULL Pointer

- Many functions return NULL to indicate an error has occurred

```
/* these are all equivalent */  
int *p = NULL;  
int *p = (int *)0;    // cast 0 to a pointer type  
int *p = (void *)0;   // automatically gets converted to the correct type
```

- NULL is considered "false" when used in a Boolean context
 - **Remember: false expressions** in C are defined to be zero or NULL
- The following two are equivalent (the second one is preferred for readability):


```
if (p) ...  
if (p != NULL) ...
```

What is Aliasing?

- **Two or more** variables are **aliases** of each other when they all reference the same memory (so different names, same memory location)
- When one pointer is copied to another pointer it *creates an alias*
- **Side effect**: Changing one variables value (content) changes the value for other variables
 - Multiple variables all read and write the same memory location
 - Aliases occur either by accident (coding errors) or deliberate (**careful: readability**)

```
int i = 5;  
int *p;  
int *q;  
p = &i;
```

```
q = p;    // *p & *q are aliases  
*q = 4;   // changes i
```



Result *p, *q and i all have the value of 4

Defining Arrays

Definition: `type name[count]`

- **"Compound"** data type where each value in an array is an element of `type`
- Allocates **name** with a *fixed* `count` array elements of type `type`
- Allocates (`count * sizeof(type)`) bytes of **contiguous memory**
- Common usage is to specify a compile-time constant for `count`

```
#define BSZ 6  
int b[BSZ];
```

BSZ is a macro replaced by the C preprocessor at compile time


- Array **names are constants** and **cannot be assigned** (the name cannot appear on the Lside by itself)

```
a = b;           // invalid does not copy the array  
                // copy arrays element by element
```

b: 0x5000

```
int b[6];
```

1 word
(int = 4 bytes)




Array Initialization

- Initialization: `type name[count] = {val0,...,valN};`
 - `{ }` (*optional*) initialization list can only be used at **time of definition**
 - If no `count` supplied, `count` is determined by compiler using the number of array initializers
- `int block[20] = {};` *//only works with constant size arrays*
 - defines an **array of 20 integers** each element filled with zeros
 - Performance comment: do not zero automatic arrays unless really needed!
- When a `count` is given:
 - extra initialization values** are **ignored**
 - missing initialization values** are set to **zero**

```
int block[5] = {2, 3, 5, 6, 11, 13};
```

not needed and if used may truncate initialization list

6 initialization values given, **only 5 are used**



Accessing Arrays Using Indexing


- `name[index]` selects the `index` element of the array
 - `index` should be unsigned
 - Elements range from: 0 to `count - 1` (`int x[count];`)
- `name[index]` can be used as an assignment target or as a value in an expression
- Array name (by itself with no `[]`) on the Rside evaluates to the address of the first element of the array

```
int a[5];  
int b[5];
```

```
int b[5];  
int *p = b;
```

p 9000

1 word
(int = 4 bytes)




How many elements are in an array?

- The number of elements of space allocated to an array (called **element count**) and indirectly the total size in bytes of an array **is not stored anywhere!!!!!!**
- An **array name** is just the **address of the first element in a block of contiguous memory**
 - So an array does not know its own size!

```
#define SZ 6
int block[SZ];      // you specify the array has SZ elements
int indx;           // use when SZ is defined

for (indx = 0; indx < SZ; indx++)
    block[indx] = 0;
```

```
int b[6];
```



Determining Element Count for a compiler calculated array

- Programmatically determining the element count in a compiler calculated array
`sizeof(array) / sizeof(of just one element in the array)`
- `sizeof(array)` only works when used in the SAME **scope** as where the array variable was **defined**

```
#include <stddef.h>

int block[] = {2, 3, 5, 6, 11, 13};    // automatic: compiler calculates array size

int cnt = (int)(sizeof(block) / sizeof(block[0])); // in this case cnt = 6

for (int indx = 0; indx < cnt; indx++)
    block[indx] = 0;
```


Pointer and Arrays - 1

- A few slides back we stated: **Array name** (by itself) on the Rside evaluates to the **address of the first element of the array**

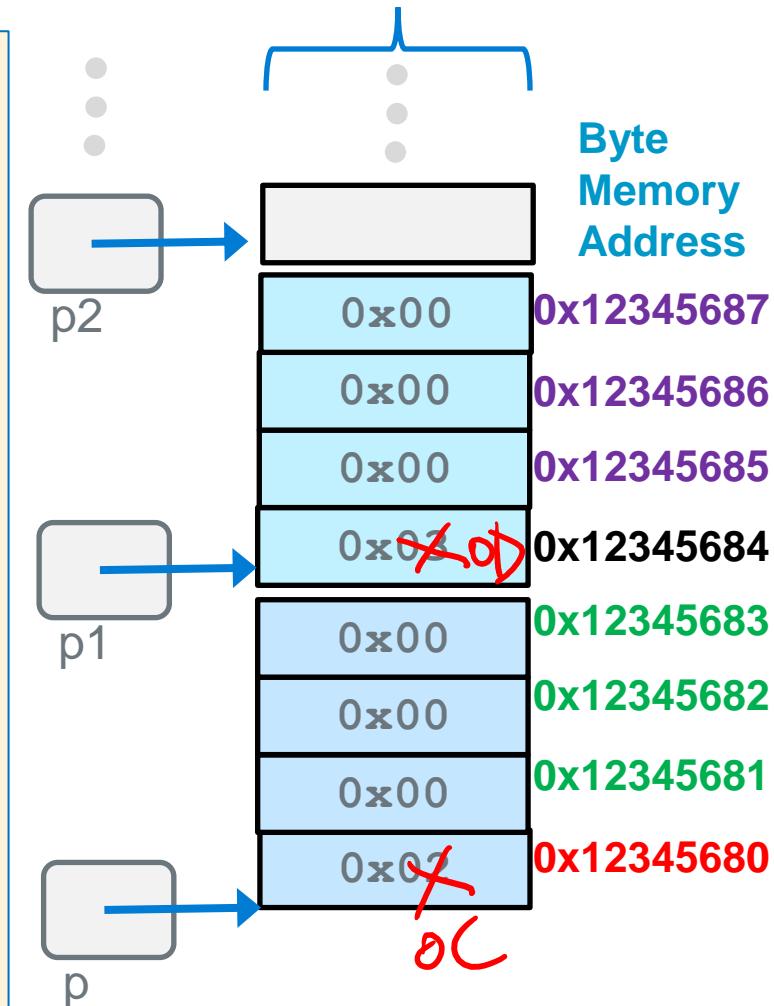
```
int buf[] = {2, 3, 5, 6, 11};
```

- Array indexing syntax (`[]`) an operator that performs **pointer arithmetic**
- buf** and **&buf[0]** on the **Rside are equivalent**, **both evaluate** to the address of the first array element

```
int *p = buf;           // or int *p = &buf[0];
int *p1 = &buf[1];
int *p2 = &buf[2];
int *p3 = &buf[3];

*p = *p + 10;
*p1 = *p1 + 10;         // {12, 13, 5, 6, 11}
```

1 byte Memory Content
One byte per row



Pointer and Arrays - 2

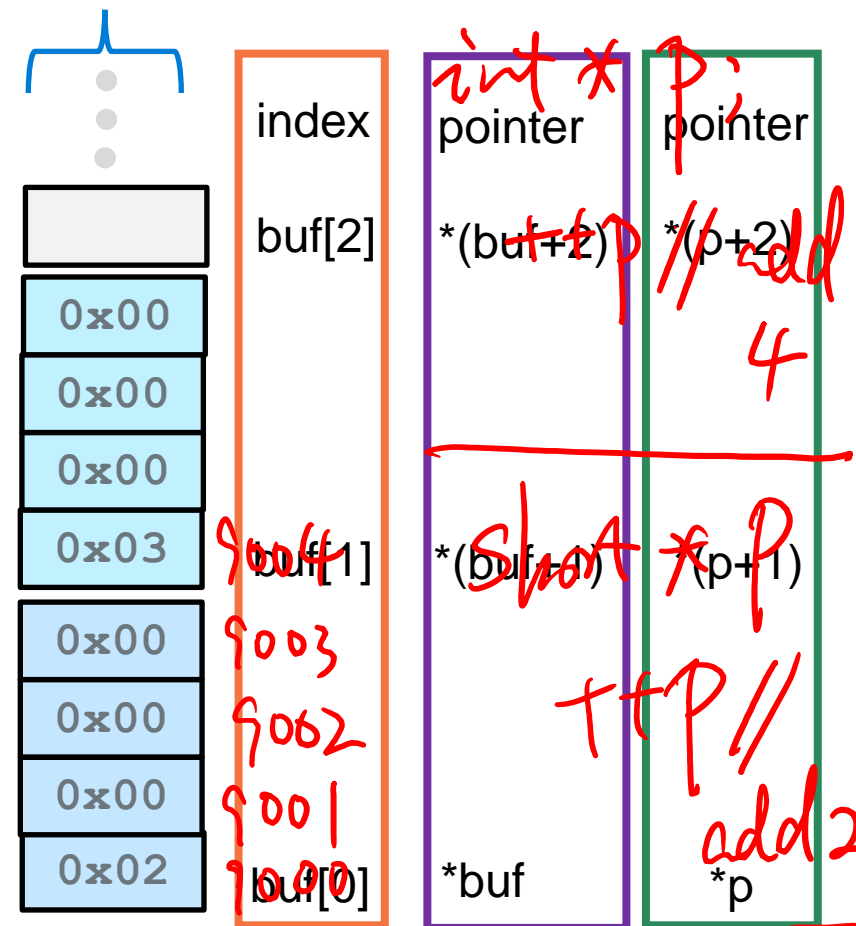
When `p` is a pointer, the actual value of `(p+1)` depends on the type that pointer `p` points at

- `(p+1)` adds `1 x sizeof(what p points at)` bytes to `p`
 - `++p` is equivalent to `p = p + 1`
- Using pointer arithmetic to find array elements:
 - Address of the second element `&buf[1]` is `(buf + 1)`
 - It can be referenced as `*(buf + 1)` or `buf[1]`

```
int buf[] = {2, 3, 5, 6, 11};
int *p;
p = buf;

*p = *p + 10;
*(p + 1) = *(p + 1) + 10; // {12, 13, 5, 6, 11}
```

1 byte Memory Content
One byte per row




Pointer Arithmetic In Use – C's Performance Focus

*char * p = a;
++p; ✓*

```
char a[] = {'A', 'B', 'C'};
```

++a ✗



a[1] ⇒ a+1

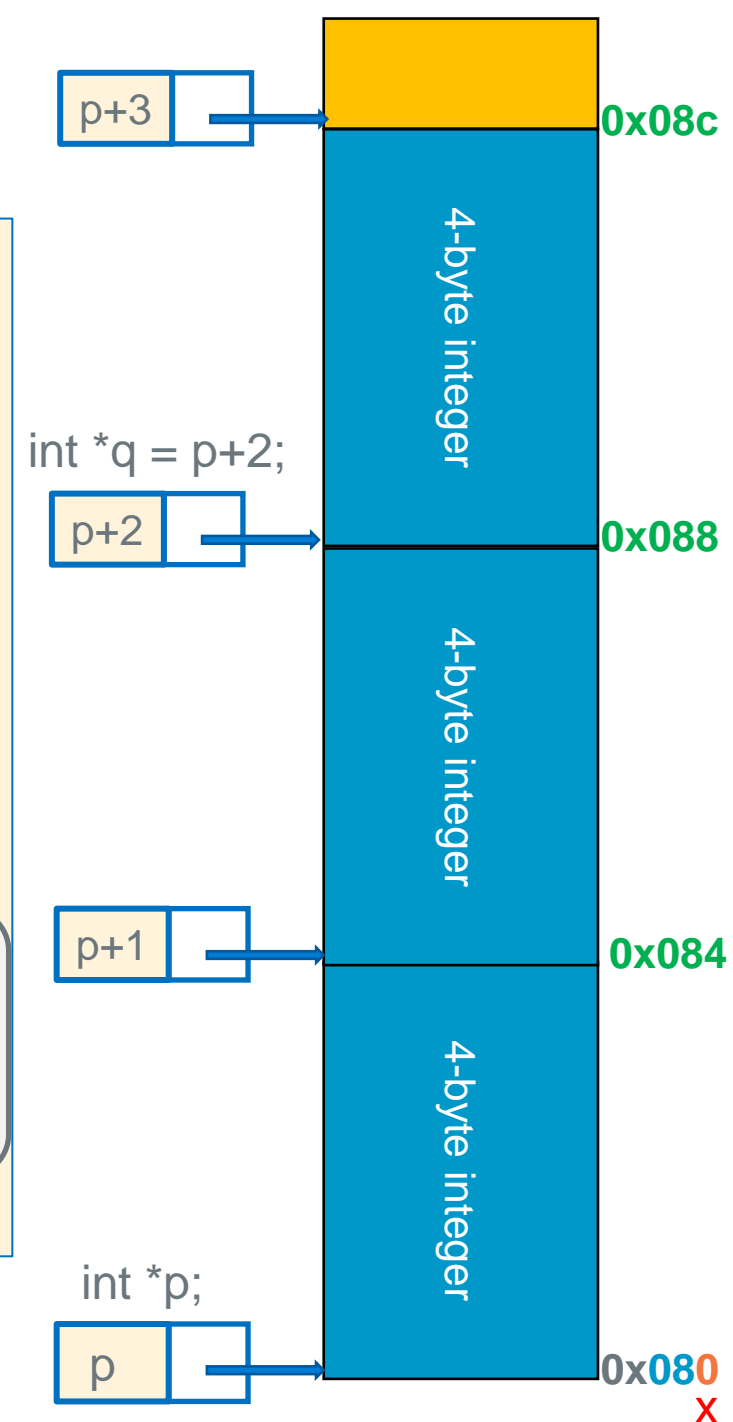
- **Alert!:** C performance focus does not perform any array “bounds checking”
- **Performance by Design:** *bound checking slows down execution of a properly written program*
- **Example:** array **a** of length **i**, C does not verify that **a[j]** or ***(a + j)** is valid (does not check: $0 \leq j < i$)
 - C simply “*translates*” and accesses the memory specified from: **a[j]** to be ***(a + j)** which may be *outside the bounds* of the array
 - OS only “**faults**” for an incorrect access to memory (read-only or not assigned to your process)
 - It does not fault for out of bound indexes or out of scope
- **lack of bound checking** is a **common source of errors and bugs** and is a common criticism of C

Pointer Arithmetic

- You cannot add two pointers (*what is the reason?*)
- A pointer *q* can be subtracted from another pointer *p* when the pointers are the same type – **best done only within arrays!**
- The value of $(p - q)$ is the number of **elements between** the two pointers
 - Using memory address arithmetic (*p* and *q* Rside are both **byte addresses**):

distance in elements = $(p - q) / \text{sizeof}(*p)$

$$(p + 3) - p = 3 = (0x08c - 0x080) / 4 = 3$$




Pointer Comparisons

- Pointers (**same type**) can be compared with the comparison operators:

<, <=, ==, !=, >=, >

```
int numb[] = {9, 8, 1, 9, 5};  
int *end;  
int *a;  
end = numb + (int) (sizeof(numb)/sizeof(*numb));  
a = numb;  
while (a < end) // compares two pointers (address)  
    /* rest of code including doing an a++ */
```



- Invalid, Undefined, or **risky** pointer arithmetic (some examples)
 - Add, multiply, divide on two pointers
 - Subtract two pointers of different types or pointing at different arrays
 - Compare two pointers of different types
 - Subtract a pointer from an integer

Fast Ways to "Walk" an Array: Use a Limit Pointer

```
int x[] = {0xd4c3b2a1, 0xd4c3b200, 0x12345684};  
int cnt = (int)(sizeof(x) / sizeof(*x));
```

```
int *ptr;  
int *xpt;  
ptr = x; //or &x[0]  
xpt = ptr + cnt;
```

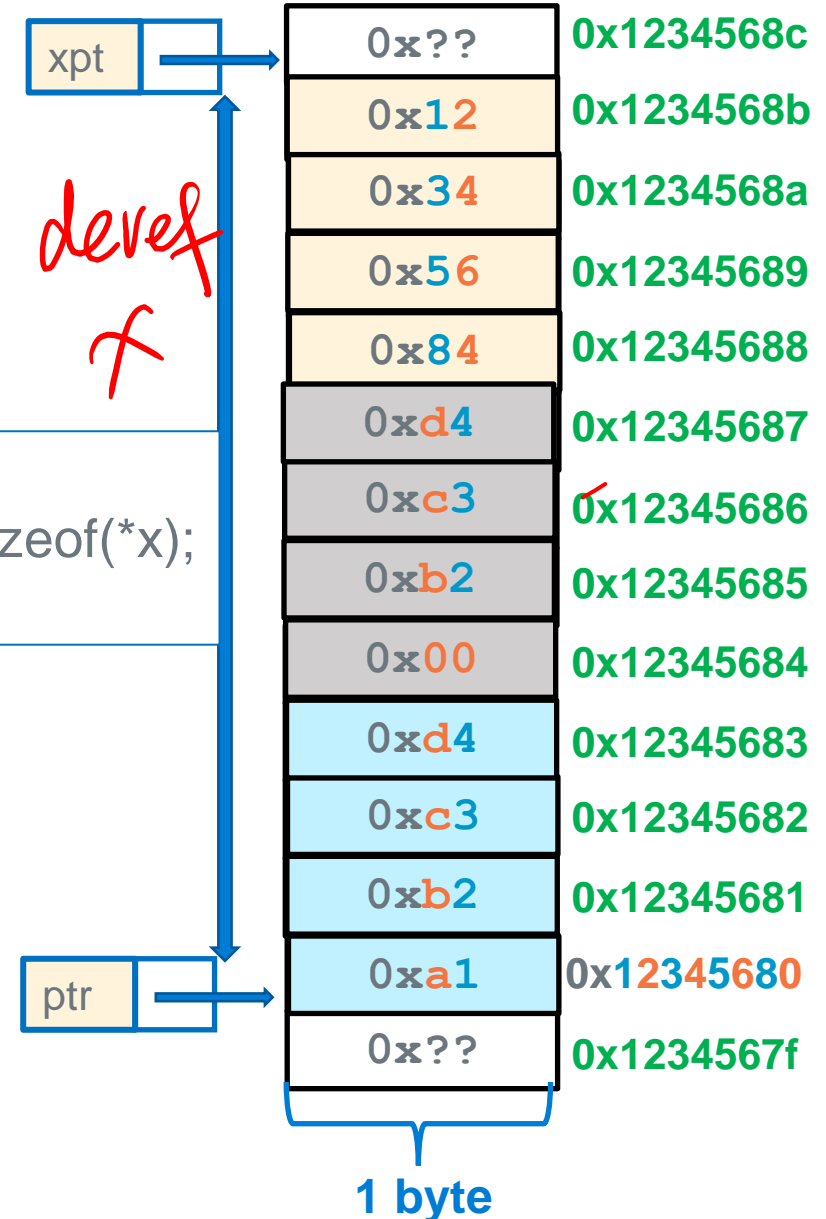
xpt is a loop **limit pointer**
it points 1 element past
the end of the array

```
while (ptr < xpt) {  
    printf("%#x\n", *ptr);  
    ptr++;  
}
```

```
% ./a.out  
0xd4c3b2a1  
0xd4c3b200  
0x12345684
```

cnt = 3;
bytes = cnt * sizeof(*x);
= 12

X X : deref



C Precedence and Pointers


- ++ -- pre and post increment combined with pointers can create code that is complex, hard to read and difficult to maintain
- Use () to help readability

Operator	Description	Associativity
() [] . -> ++ --	Parentheses or function call Brackets or array subscript Dot or Member selection operator Arrow operator Postfix increment/decrement	left to right
++ -- + - ! ~ (type) * & sizeof	Prefix increment/decrement Unary plus and minus not operator and bitwise complement type cast Indirection or dereference operator Address of operator Determine size in bytes	right to left
* / %	Multiplication, division and modulus	left to right
+ -	Addition and subtraction	left to right
<< >>	Bitwise left shift and right shift	left to right
< <= > >=	relational less than/less than equal to relational greater than/greater than or equal to	left to right
== !=	Relational equal to or not equal to	left to right
&&	Bitwise AND	left to right
^	Bitwise exclusive OR	left to right
	Bitwise inclusive OR	left to right
&&	Logical AND	left to right
	Logical OR	left to right
? :	Ternary operator	right to left
= += -= *= /= %= &= ^= = <<= >>=	Assignment operator Addition/subtraction assignment Multiplication/division assignment Modulus and bitwise assignment Bitwise exclusive/inclusive OR assignment	right to left
,	comma operator	left to right

common	With Parentheses	Meaning
*p++	*(p++)	(1)The Rvalue is the object that p points at (2)increment pointer p to next element ++ is higher than *
(*p)++		(1)Rvalue is the object that p points at (2)increment the object
*++p	*(++p)	(1)Increment pointer p first to the next element (2)Rvalue is the object that the incremented pointer points at
++*p	++(*p)	Rvalue is the incremented value of the object that p points at


C Strings - 1

- C does not have a **dedicated type** for strings
- Strings are an **array of characters** terminated by a **sentinel termination character**
- `'\0'` is the **Null termination character**; has the **value of zero** (do not confuse with `'0'`)
- An **array of chars** contains **a string only when** it is terminated by a `'\0'`
- **Length of a string** is the number of characters in it, not including the `'\0'`
- Strings in C are not objects
 - **No embedded information about them**, you just have a **name** and a memory **location**
 - You **cannot use** `+` or `+=` to concatenate strings in C
 - For example, you must **calculate string length** using code at runtime looking for the end



C Strings - 2

- First '`\0`' encountered from the start of the string always indicates the end of a string
- The '`\0`' **does not have to be** in the **last element in the space allocated to the array**
 - But, String length is always **less than the size of the array** it is contained in
- In the example below, the array `buf` contains two strings
 - One string starts at `&(buf[0])` is "`cat`" with a string length of 3
 - The other string starts at `&(b[4])` is "`o`" with a string length of 1
 - "`o`" has two bytes: '`o`' and '`\0`'



Defining Strings: Initialization

- When you combine the automatic length definition for arrays with double quote(") **initialization**
 - Compiler automatically adds the null terminator '\0' for you

```
char a[4] = {'c', 'a', 't', '\0'};  
char b[] = "cat";  
char c[] = {'c', 'a', 't', '\0', 'a', 'b'};  
char empty[] = "";
```

// compiler calculates size, adds '\0'
// array size 6, string length 3
// empty string - contains '\0'
// string length = 0

Defining Strings: Initialization Equivalents

- Following definitions create **equivalent** 4-character arrays
 - These are all strings as they all include a null ('\0') terminator

```
char a[4] = {'c', 'a', 't', '\0'};
char b[4] = {'c', 'a', 't', 0};
char c[4] = {'c', 'a', 't'};           // missing initial value defaults to 0
char d[4] = { 99, 97, 116, 0};         // 99 = 'c', 97 = 'a', 116 = 't'
char e[4] = "cat";
char f[4] = "cat\0";                   // literal has 5 chars; array f string
                                       // length is 3
```

Background: Different Ways to Pass Parameters

- **Call-by-reference (or pass by reference)**

- Parameter in the called function is an alias (references the same memory location) for the supplied argument
- Modifying the parameter modifies the calling argument

Call-by-value (or pass by value) (C)


- What **Called** Function Does
 - Passed Parameters are used like local variables
 - Modifying the passed parameter in the function is allowed just like a local variable
 - So, writing to the parameter, only changes the copy
- The return value from a function in C is **by value**

Passing Parameters – Call by Value Example

```
int main(void)
{
    int x = 5;
    inc(x); // makes a copy of x
    printf("%d\n", x); // 5 or 6 ?
}

void inc(int i) // i is local to inc
{
    ++i;
}
```

if this was an expression like `inc(x+1)` it evaluates and stores the result in the memory allocated for the copy



- when `inc(x)` is called, a copy of `x` is made to another memory location
 - `inc()` cannot change the variable `x` since `inc()` does not have the address of `x`, it is local to `main()` so, 5 is printed
- The `inc()` function is free to change its copy of the argument (just like any local variable) remember it does NOT change the parameter in `main()`

Output Parameters

- Passing a pointer parameter with the intent that the called function will use the address it to store values for use by the calling function, then pointer parameter is called an **output parameter**
- To pass the address of a variable `x` use the **address operator** (`&x`) or the contents of a pointer variable that points at `x`, or the name of an array (the arrays address)
- To be receive an address in the called function, define the corresponding parameter type to be a pointer (add `*`)
 - It is common to describe this method as: “pass a pointer to `x`”
- C is still using “*pass by value*”
 - we pass the **value** of the address/pointer in a **parameter copy**
 - **The called routine** uses the address to change a variable in the caller's scope

```
void inc(int *p);  
int  
main(void)  
{  
    int x = 5;  
    inc(&x);  
    printf("%d\n", x);  
    return EXIT_SUCCESS;  
}  
  
void  
inc(int *p)  
{  
    if (p != NULL)  
        *p += 1; // or (*p)++  
}
```

Example Using Output Parameters

Pass the address of x (&x)

```
void inc(int *p);
int
main(void)
{
    int x = 5;
    inc(&x);
    printf("%d\n", x);
    return EXIT_SUCCESS;
}
```


Receive an address copy (int *p)

```
void
inc(int *p)
{
    if (p != NULL)
        *p += 1;    // or (*p)++
}
```

Write to the output variable (*p)

At the Call to inc() in main()

1. Allocate space for p
2. Copy x's address into p



With a pointer to X,

inc() can change x in main()

this is called a side effect

p just like any other local variable

Array Parameters: Call-By-Value or Call-By-Reference?

- `Type []` array parameter is automatically “**promoted**” to a pointer of type `Type *`, and a copy of the *pointer* is *passed by value*

```
int main(void)
{
    int numbers[] = {9, 8, 1, 9, 5};

    passa(numbers);
    printf("numbers size:%lu\n", sizeof(numbers)); // 20
    return EXIT_SUCCESS;
}
```

```
void passa(int a[])
{
    printf("a size:%lu\n", sizeof(a)); // 4
    return;
}
```

IMPORTANT:

See the size difference 20 in `main()` in `passa()` is 4 bytes (size of a pointer)

- Call-by-value pointer (callee can change the pointer parameter to point to something else!)
- Acts like **call-by-reference** (called function can change the contents caller's array)

Arrays As Parameters: What is the size of the array?

- It's tricky to use arrays as parameters, as **they are passed as pointers to the start of the array**
 - In C, Arrays do not know their own size and at runtime there is no “bounds” checking on indexes

```
int sumAll(int a[]);  
  
int main(void)  
{  
    int numb[] = {9, 8, 1, 9, 5};  
    int sum = sumAll(numb);  
  
    return EXIT_SUCCESS;  
}  
  
int sumAll(int a[])  
{  
    int i, sum = 0;  
    int sz = (int) (sizeof(a)/sizeof(*a));  
    for (i = 0; i < sz; i++) // this does not work  
        sum += a[i];  
}
```

the name is the address, so this is passing a pointer to the start of the array

“inside” the body of sumAll(), the question is: how big is that array? all I have is a POINTER to the first element.....
sz is a 1 on 32 bit arm

Arrays As Parameters, Approach 1: Pass the size

Two ways to pass array size


1. pass the **count** as an additional argument
2. add a **sentinel element** as the last element

remember you can only use `sizeof()` to calculate element count where the array is defined

```
int sumAll(int *a, int size);
int main(void)
{
    int numb[] = {9, 8, 1, 9, 5};
    int cnt = sizeof(numb)/sizeof(numb[0]);

    printf("sum is: %d\n", sumAll(numb, cnt));
    return EXIT_SUCCESS;
}
```

1 word content
(int = 4 bytes)



```
int sumAll(int *a, int size)
{
    int sum = 0;
    int *end;
    end = a + size;

    while (a < end)
        sum += *a++;
    return sum;
}
```

same as:
sum = sum + *a;
a++;

Arrays As Parameters, Approach 2: Use a sentinel element

- A **sentinel** is an element that contains a value that is not part of the normal data range
 - Forms of 0 are often used (like with strings). Examples: '\0', NULL

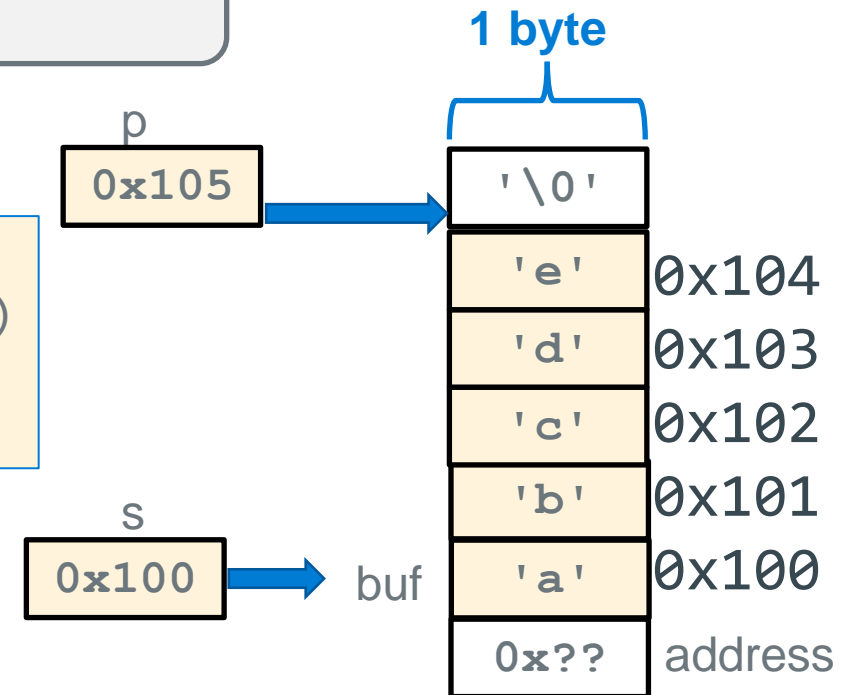
```
int strlen(char *a);
int main(void)
{
    char buf[] = {'a', 'b', 'c', 'd', 'e', '\0'}; // string

    printf("Number of chars is: %d\n", strlen(buf));
    return EXIT_SUCCESS;
}
```

/ Assumes parameter is a terminated string */*

```
int strlen(char *s)
{
    char *p = s;
    if (p == NULL)
        return 0;
    while (*p++)
        ;
    return (p - s - 1);
}
```

same as:
while (*p != '\0')
 p = p + 1;
return (p - s);



Do not overuse strlen()

- C string library function `strlen()` calculates string length **at runtime**
- **Do not overuse strlen(), as it walks the array each time called**

```
int count_e(char *s) //  $O(n^2)$  !!!
{
    int count = 0;
    if (s == NULL)
        return 0;
    for (int j = 0; j < strlen(s); j++) {
        if (s[j] == 'e')
            count++;
    }
    return count ;
}
```

```
int count_e(char *s) //  $O(n)$  !!!
{
    int count = 0;
    if (s == NULL)
        return 0;
    while (*s) {
        if (*s++ == 'e')
            count++;
    }
    return count ;
}
```

same as:

```
while (*s != '\0') {
    if (*s == 'e')
        count++;
    s++;
}
```

Comparing strings

- Characters can be easily compared ($c1 < c2$) as they are numbers, so the **character order** is determined by the ASCII values assigned to each character
 - 65 = A 66 = B 67 = C 68 = D 69 = E 70 = F 71 = G, and so on.
- **Example:** the following strings are in lexicographical (alphabetical) order:
"" "a" "az" "c" "cab" "cabin" "cat" "catastrophe"
- Compare two strings lexicographically (i.e., comparing ASCII values), subtract one from the other

Return Value	Comparison
< 0	$s1 < s2$
> 0	$s1 > s2$
$= 0$	$s1 == s2$

```
int strcmp(char *s1, char *s2)
{
    while (*s1 == *s2) {
        if ((*s1 == '\0') || (*s2 == '\0'))
            break;
        s1++;
        s2++;
    }
    return *s1 - *s2; // character difference
}
```

Copying Strings: Use the Sentinel; libc: strcpy(), strncpy()

- To copy an array, you must copy each character from source to destination array
- Watch overwrites: strcpy assumes the target array size is equal or larger than source array

index	0	1	2	3	4	5
char	'H'	'e'	'l'	'l'	'o'	'\0'

```
char str1[80];  
strcpy(str1, "hello");
```

```
char *strcpy(char *s0, char *s1)  
{  
    char *str = s0;  
  
    if ((s0 == NULL) || (s1 == NULL))  
        return NULL;  
    while (*s0++ = *s1++)  
        ;  
    return str;  
}
```

```
// strncpy adds a length limit on copy  
char str1[6];  
strncpy(str1, "hello", 5); // \0 not copied  
str1[5] = '\0'; // make sure \0 terminated
```

```
char *strncpy(char *s0, char *s1, int len)  
{  
    char *str = s0;  
    if ((s0 == NULL) || (s1 == NULL))  
        return NULL;  
  
    while ((*s0++ = *s1++) && --len)  
        ;  
    return str;  
}
```

Reference: Some String Routines in libc (#include <string.h>)

Function	Description
<code>strlen(<i>str</i>)</code>	returns the # of chars in a C string (before null-terminating character).
<code>strcmp(<i>str1</i>, <i>str2</i>)</code> , <code>strncmp(<i>str1</i>, <i>str2</i>, <i>n</i>)</code>	compares two strings; returns 0 if identical, <0 if <i>str1</i> comes before <i>str2</i> in alphabet, >0 if <i>str1</i> comes after <i>str2</i> in alphabet. <i>strncmp</i> stops comparing after at most <i>n</i> characters.
<code>strchr(<i>str</i>, <i>ch</i>)</code> <code>strrchr(<i>str</i>, <i>ch</i>)</code>	character search: returns a pointer to the first occurrence of <i>ch</i> in <i>str</i> , or <i>NULL</i> if <i>ch</i> was not found in <i>str</i> . <i>strrchr</i> find the last occurrence.
<code>strstr(<i>haystack</i>, <i>needle</i>)</code>	string search: returns a pointer to the start of the first occurrence of <i>needle</i> in <i>haystack</i> , or <i>NULL</i> if <i>needle</i> was not found in <i>haystack</i> .
<code>strcpy(<i>dst</i>, <i>src</i>)</code> , <code>strncpy(<i>dst</i>, <i>src</i>, <i>n</i>)</code>	copies characters in <i>src</i> to <i>dst</i> , including null-terminating character. Assumes enough space in <i>dst</i> . Strings must not overlap. <i>strncpy</i> stops after at most <i>n</i> chars, and <u>does not</u> add null-terminating char.
<code>strcat(<i>dst</i>, <i>src</i>)</code> , <code>strncat(<i>dst</i>, <i>src</i>, <i>n</i>)</code>	concatenate <i>src</i> onto the end of <i>dst</i> . <i>strncat</i> stops concatenating after at most <i>n</i> characters. <u>Always</u> adds a null-terminating character.
<code>strspn(<i>str</i>, <i>accept</i>)</code> , <code>strcspn(<i>str</i>, <i>reject</i>)</code>	<i>strspn</i> returns the length of the initial part of <i>str</i> which contains <u>only</u> characters in <i>accept</i> . <i>strcspn</i> returns the length of the initial part of <i>str</i> which does <u>not</u> contain any characters in <i>reject</i> .

2D Array of Char (where elements may contain strings)

- 2D array of chars (where rows may include strings)
- Each row has the same fixed number of memory allocated
- All the rows are the same length regardless of the actual string length)
- The column size must be large enough for the longest string

high memory char aos2d[3][22] = {"my", "two dimensional", "char array"};

aos2d[2]	c	h	a	r		a	r	r	a	y	'\0'											
aos2d[1]	t	w	o		d	i	m	e	n	s	i	o	n	a	l		a	r	r	a	y	'\0'
aos2d[0]	m	y	'\0'																			

low memory

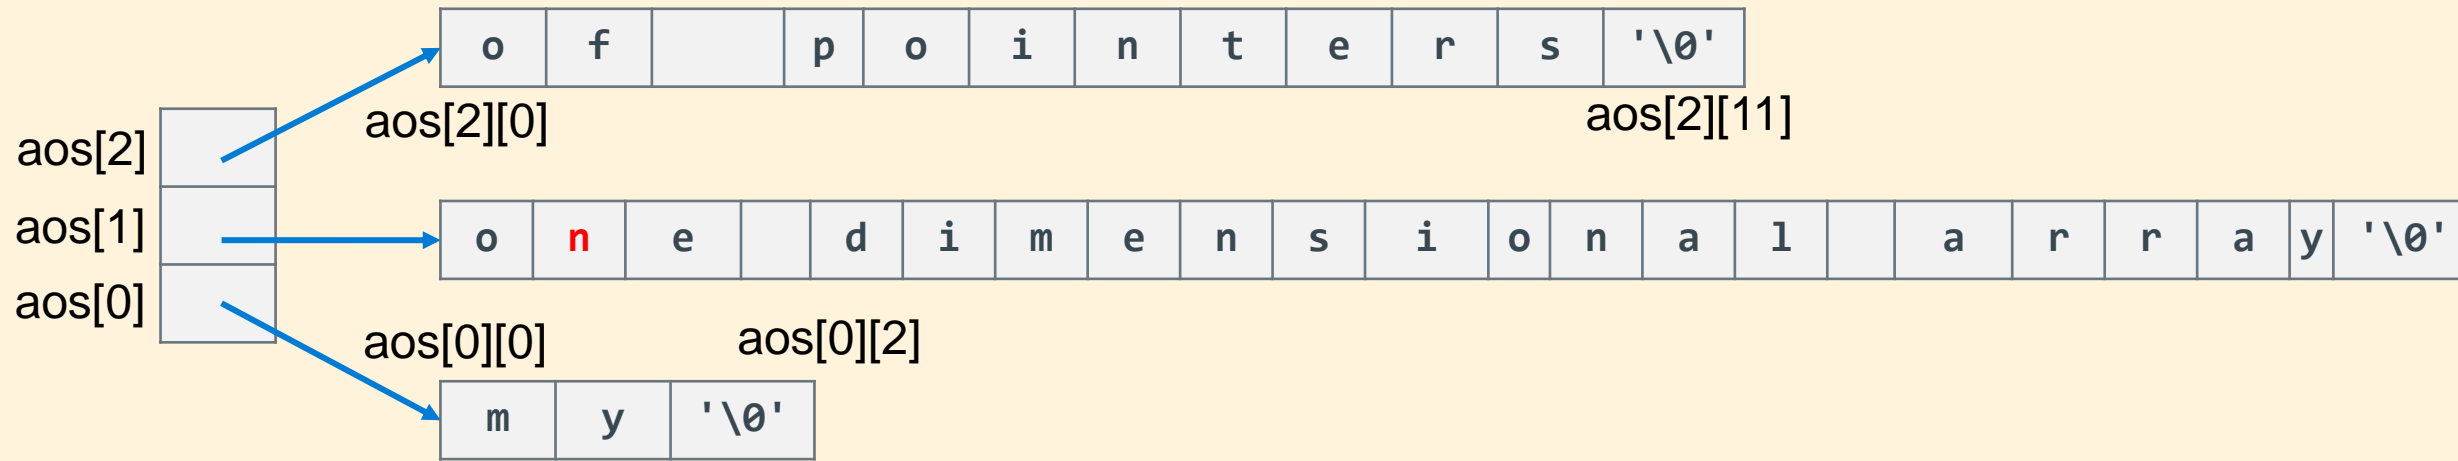
```
#define ROWS 3
char aos[ROWS][22] = { "my", "two dimensional", "char array"};
char (*ptc)[22] = aos; // ptr points at a row of 22 chars
```

```
for (int i = 0; i < ROWS; i++)
    printf("%s\n", *(ptc + i));
```

high memory

Pointer Array to Strings (This is NOT a 2D array)

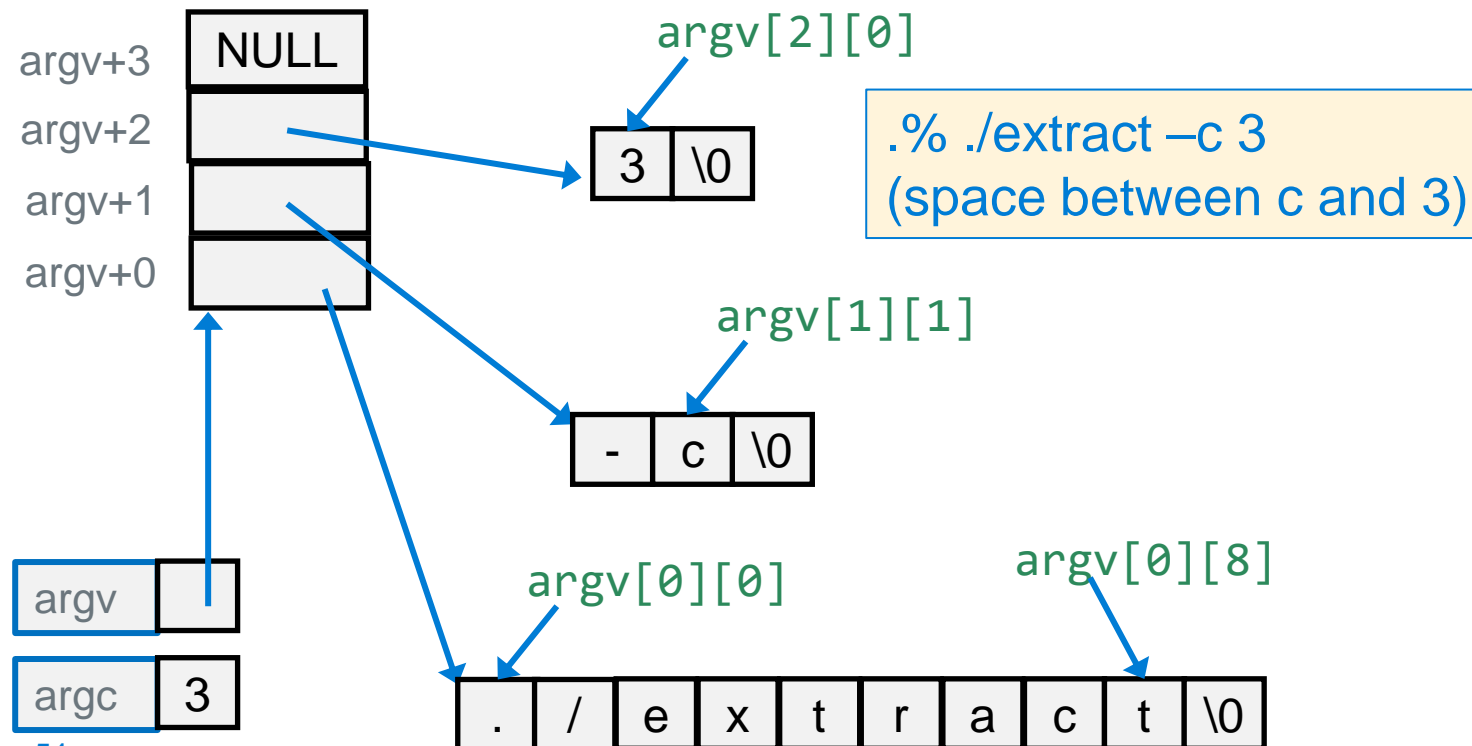
- 2D char arrays are an inefficient way to store strings (wastes memory) unless all the strings are similar lengths, so 2D char arrays are rarely used with string elements
- An array of pointers is common for strings as "rows" can vary in length



- aos is an array of pointers; each pointer points at a character array (also a string here)
- Not a 2D array, but any char can be accessed as if it was in a 2D array of chars

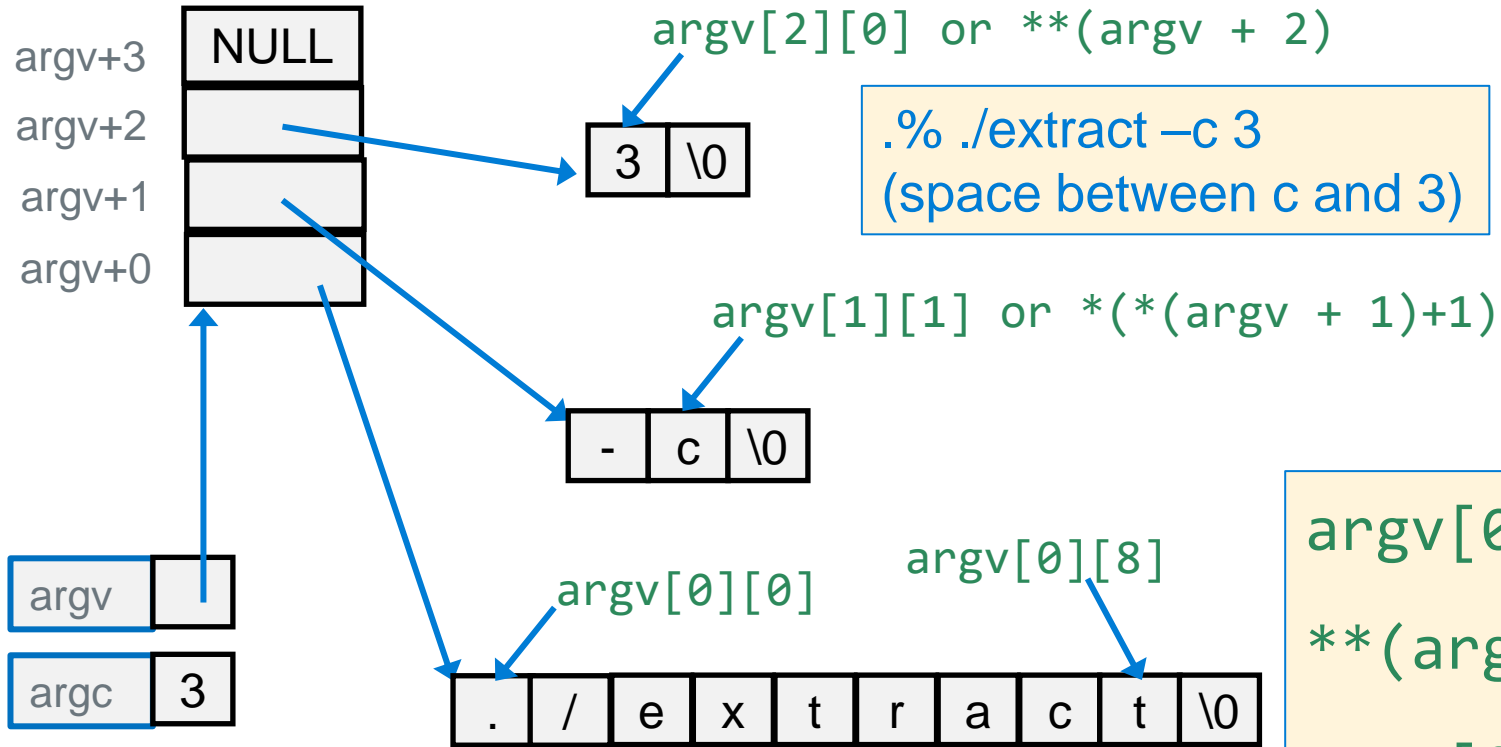
main() Command line arguments: argc, argv

- Arguments are passed to main() as a pointer to an array of pointers (****argv** or ***argv[]**)
Conceptually: % ***argv[0]** ***argv[1]** ***argv[2]**
- argc** is the number of VALID elements (they point at something)
- *argv** (**argv[0]**) is **usually** is the **name** of the executable file (% **./vim** file.c)
- *(argv + argc)** always contains a NULL (0) sentinel
- *argv[]** (or ****argv**) elements point at **mutable strings!**



```
printf("%s\n", *(argv+0));  
printf("%s\n", *(argv+1));  
printf("%s\n", *(argv+2));
```

main() Command line arguments: argc, argv



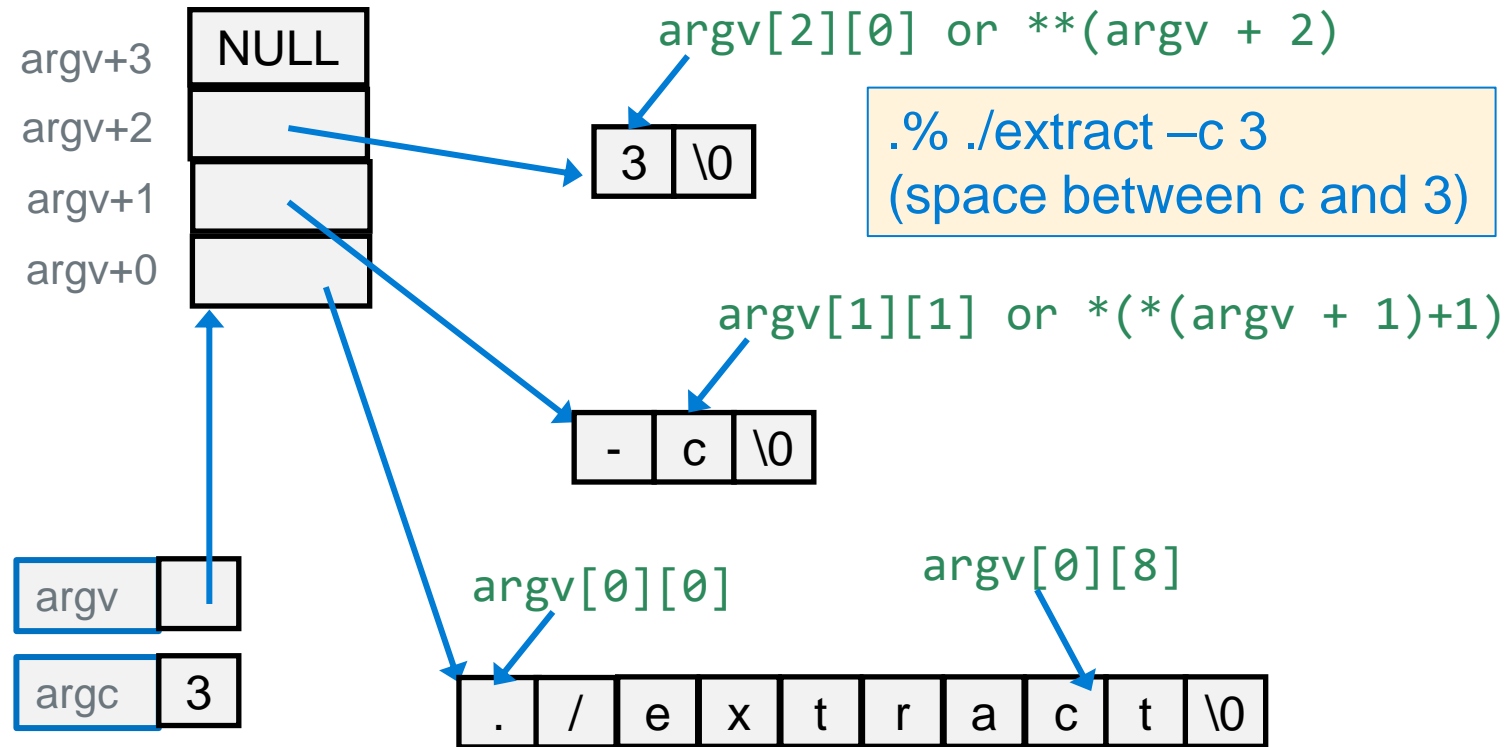
`argv[0][0]` equiv to `**argv`
`**argv` equiv `**argv`
`argv[0][8]` equiv `*(argv + 8)`

`char *pt = *argv;`

`*pt` equiv to `**argv`

`*(pt+8)` equiv to `*(argv + 8)`

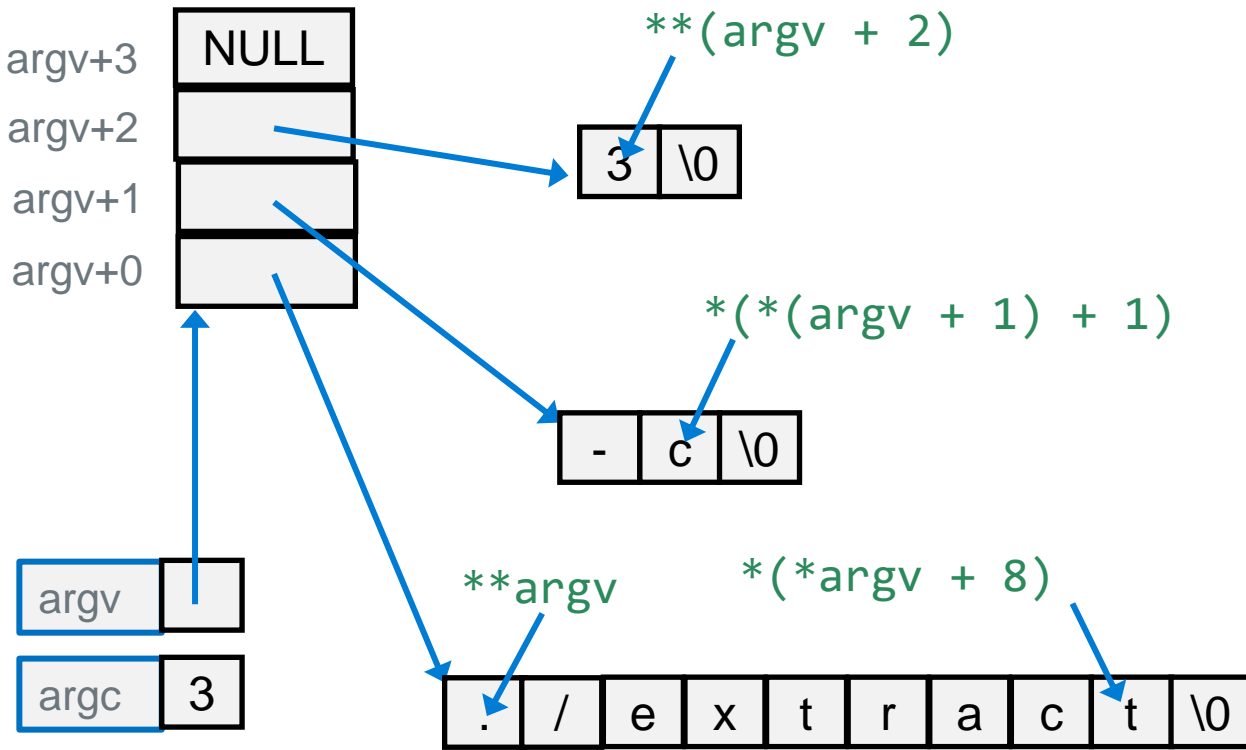
main() Command line arguments: argc, argv



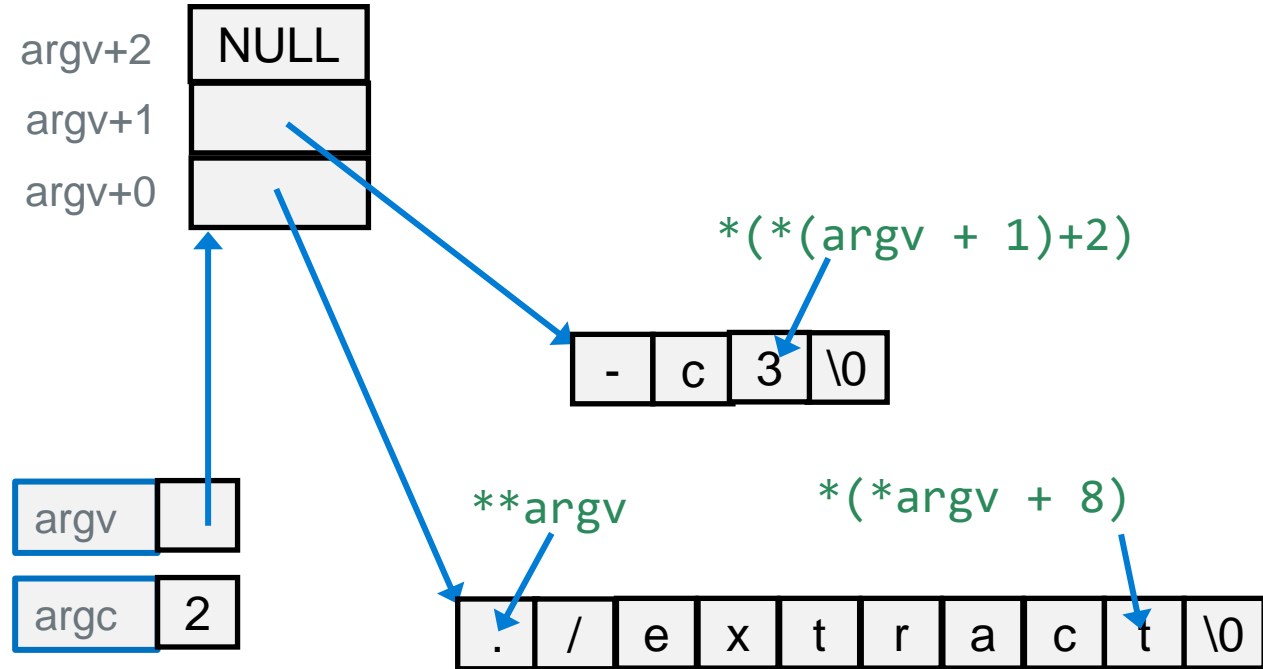
```
int main(int argc, char *argv[])
{
    for (int i = 0; argv[i] != NULL; i++) {
        for (int j = 0; argv[i][j] != '\0'; j++)
            putchar(argv[i][j]);
        putchar('\n');
    }
    return EXIT_SUCCESS;
}
```

```
int main(int argc, char **argv)
{
    char *pt;
    while ((pt = *argv++) != NULL) {
        while (*pt != '\0')
            putchar(*pt++);
        putchar('\n');
    }
    return EXIT_SUCCESS;
}
```

main() Command line arguments: argc, argv



`./extract -c 3`
(space between c and 3)



`./extract -c3`
(No space between c and 3)

PA4: Creating a 2D Array of Mutable String Pointers


1. Break a string of comma separated words into individual strings without copying. Do This by walking the string until you see an either a comma , or a newline \n. Each points at a field or column in a record.
2. Record the start of each string into successive elements in an array of pointers
3. Replace each comma or newline with a null '\0'

char *buf



buf[0]	buf[1]	buf[2]	buf[3]	buf[4]	buf[5]	buf[6]	buf[7]	buf[8]	buf[9]	buf[10]	buf[11]	buf[12]	buf[13]
c	s	e	\0	1	0	0	\0	L	i	n	e	\0	\0

char **ptable



ptable	ptable+1	ptable+2

./extract -c3


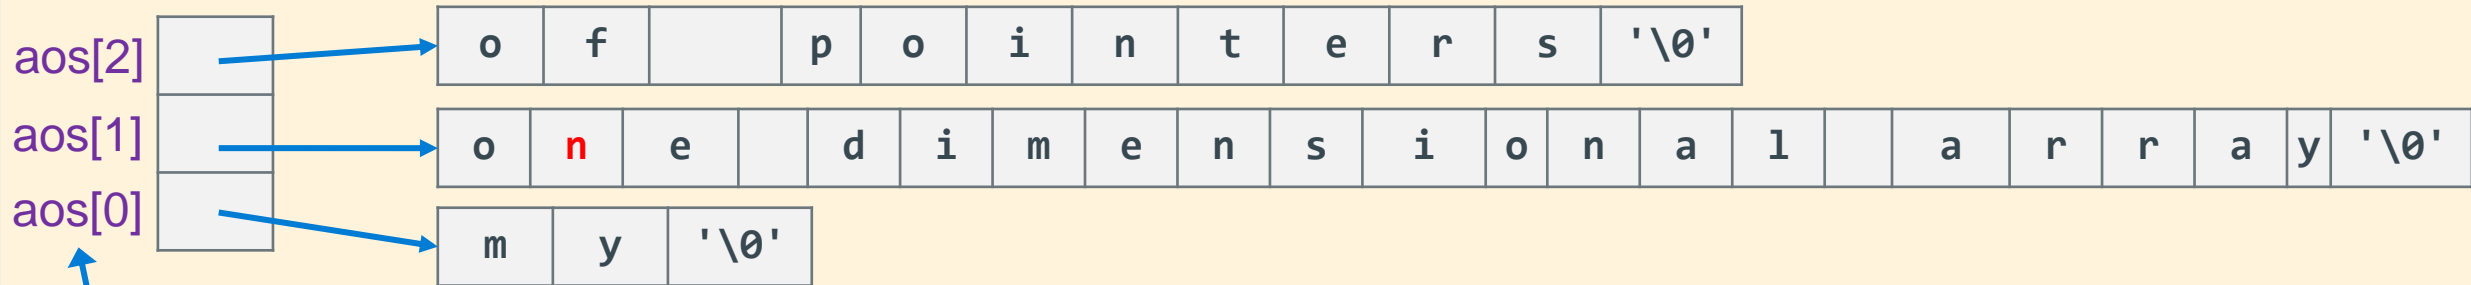
```
// extract of token(), passed buf, ptable  
and cnt
```

```
char **endptr = ptable + cnt;  
*ptable = buf;  
while ((ptable < endptr) && (*buf != '\0'))  
{  
    *ptable++ = buf;  
    while (*buf != '\0') {  
        /* process chars including buf++ */  
    }  
}  
// check for too many or too few fields
```

Review: Pointer Array to Strings

How to access: `aos[1][1]` is `*(*(aos + 1) + 1)` which contains 'n'
its address is `(*(aos + 1) + 1)`

`aos+2` is not shown due to space limits on the slide



Notice that the first elements address is the array name

```
char *ptr;  
ptr = *aos;  
*ptr = 'X';  
if (*ptr == ',')  
    or  
if (**aos) == ','
```