

UCSD CSE 30

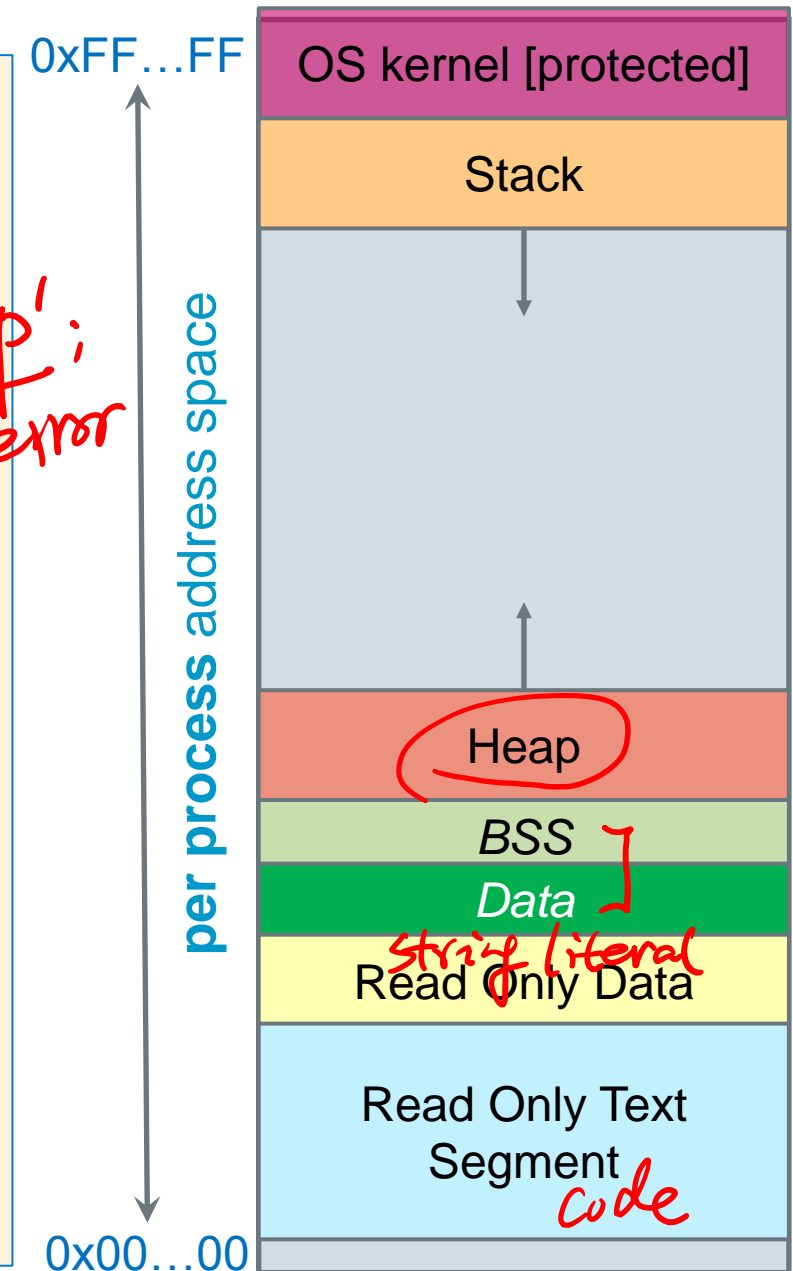
Computer Organization and Systems Programming

C Part 3

Process Memory Under Linux

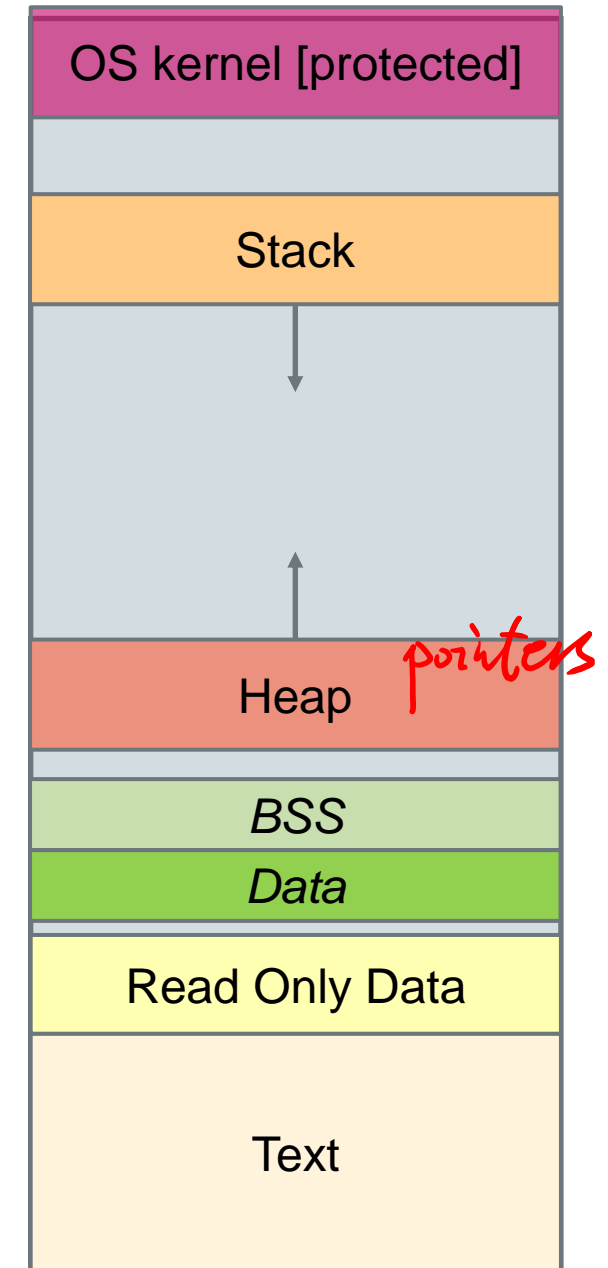
- When your **program is running** it has been **loaded into memory** and is **called a process**
- **Stack segment:** Stores **Local** variables
 - Allocated and freed at function call entry & exit
- **Data segment + BSS:** Stores **Global** and **static** variables
 - **Allocated/freed** when the process **starts/exits**
 - **BSS** - Static variables with an implicit initial value
 - **Static Data** - Initialized with an explicit initial value
- **Heap segment:** Stores **dynamically-allocated** variables
 - Allocated with a function call
 - Managed by the stdio library malloc() routines
- **Read Only Data:** Stores **immutable** Literals
- **Text:** Stores your code in machine language + libraries

*char *str = "hello";*
**str = 'p';*
error



The Heap Memory Segment

- **Heap**: "pool" of memory that is available to a program
 - Managed by C runtime library and linked to your code; **not managed by the OS**
- Heap memory is **dynamically** *"borrowed"* or *"allocated"* by calling a library function
- When heap memory is no longer needed, it is *"returned"* or *deallocated* for **reuse**
- Heap memory has a lifetime from allocation until it is deallocated
 - Lifetime is independent of the scope it is allocated in (it is like a static variable)
- If too much memory has already been allocated, the library will attempt to borrow additional memory from the OS and will fail, returning a NULL



Heap Dynamic Memory Allocation Library Functions

#include <stdlib.h>	args	Clears memory
<code>void *malloc(...)</code> <i>any kind of pointer</i>	<code>size_t size</code>	no
<code>void *calloc(...)</code>	<code>size_t nmemb, size_t memsize</code>	yes
<code>void free(...)</code>	<code>void *ptr</code>	no

- **void *** means these library functions return a pointer to **generic (untyped) memory**
 - Be careful with void * pointers and pointer math as void * points at untyped memory
 - The assignment to a typed pointer *"converts"* it from a void *
- **size_t** is an **unsigned integer data type**, the result of a **sizeof()** operator

```
int *ptr = malloc(sizeof(*ptr) * 100); // allocate an array of 100 ints
```

- **please read: % man 3 malloc**

Use of Malloc

```
void *malloc(size_t size)
```

- Returns a pointer to a **contiguous block** of `size` bytes of **uninitialized memory** from the heap
 - The block is **aligned to an 8-byte (arm32) or 16-byte (64-bit arm/intel) boundary**
 - returns **NULL** if allocation failed (also sets `errno`) **always CHECK for NULL RETURN!**
- Blocks returned on different calls to `malloc()` are not necessarily adjacent
- `void *` is implicitly cast into any pointer type on assignment to a pointer variable

```
char *bufptr;
```

int cnt;

```
/* ALWAYS CHECK THE RETURN VALUE FROM MALLOC!!!! */
```

```
if ((bufptr = (char*) malloc(cnt * sizeof(*bufptr))) == NULL) {  
    fprintf(stderr, "Unable to malloc memory");  
    return NULL;  
}
```

Using and Freeing Heap Memory

- `void free(void *p)`
 - Deallocates the whole block pointed to by `p` to the pool of available memory
 - Freed memory is used in future allocation (**expect the contents to change after freed**)
 - ~~✗~~ Pointer `p` must be the same address as **originally returned** by one of the heap allocation routines `malloc()`, `calloc()`, `realloc()`
 - Pointer argument to `free()` is not changed by the call to `free()`
- Defensive programming: **set the pointer to NULL** after passing it to `free()`

```
free(bufptr);  
bufptr = NULL;
```

```
// returns memory to the heap  
// set bufptr to NULL
```

Heap Memory "Leaks"

- A **memory leak** is when you **allocate memory** on the heap, **but never free it**

```
void  
leaky_memory (void)  
{  
    char *bytes = malloc(BLKSZ * sizeof(*bytes));  
    ...  
    /* code that never deallocates the memory */  
    return; // you lose the address in bytes when leaving scope  
}
```

- **Best practice:** free up memory **you allocated** when you no longer need it
 - If you keep allocating memory, you may run out of memory in the heap!
- **Memory leaks** may cause **long running programs to fault** when they **exhaust OS memory limits**
- **Valgrind** is a tool for finding memory leaks (not pre-installed in all linux distributions though!)

```

1 #define SZ 50
2 #include <stdlib.h>
3 int main(void)
4 {
5     char *buf;
6     if ((buf = malloc(SZ * sizeof(*buf))) == NULL)
7         return EXIT_FAILURE;
8     *(buf + SZ) = 'A';
9     return EXIT_SUCCESS;
10 }

```

Valgrind – Finding Buffer Overflows and Memory leaks

```

% valgrind -q --leak-check=full --leak-resolution=med -s ./valgexample
==651== Invalid write of size 1
==651==    at 0x10444: main (valg.c:8)
==651== Address 0x49d305a is 0 bytes after a block of size 50 alloc'd
==651==    at 0x484A760: malloc (vg_replace_malloc.c:381)
==651==    by 0x1041B: main (valg.c:6)
==651==
==651== 50 bytes in 1 blocks are definitely lost in loss record 1 of 1
==651==    at 0x484A760: malloc (vg_replace_malloc.c:381)
==651==    by 0x1041B: main (valg.c:6)
==651==
==651== ERROR SUMMARY: 2 errors from 2 contexts (suppressed: 0 from 0)

```

Writing outside of allocated buffer space

Memory not freed

More Dangling Pointers: Reusing "freed" memory

- When a pointer points to a memory location that is no longer “valid”
- Really hard to debug as the use of the return pointers may not generate a seg fault

```
char *dangling_freed_heap(void)
{
    char *buff = malloc(BLKSZ * sizeof(*buff));
    ...
    free(buff);
    return buff;
}
```

*char c = 'a';
return &c;*

- `dangling_freed_heap()` type code often causes the allocators (`malloc()` and friends) to **seg fault**
 - Because it corrupts data structures the heap code uses to manage the memory pool

strdup(): Allocate Space and Copy a String

```
char *strdup(char *s);
```

- **strdup** is a function that returns a **null-terminated**, heap-allocated string copy of the provided text
- Alternative: **malloc** and copy the string

```
char *str = strdup("Hello, world!");
```

```
*str = 'h';
```

```
free(str);
```

```
str = NULL;
```

Calloc()

```
void *calloc(size_t elementCnt, size_t elementSize)
```

calloc() variant of malloc() but zeros out every byte of memory before returning a pointer to it (so this has a runtime cost!)

- First parameter is the number of elements you would like to allocate space for
- Second parameter is the size of each element

```
// allocate 10-element array of pointers to char, zero filled  
char **arr;  
arr = calloc(10, sizeof(*arr));  
if (arr == NULL)  
    // handle the error
```

- Originally designed to allocate arrays but works for any memory allocation
 - calloc() multiplies the two parameters together for the total size
- calloc() is more expensive at runtime (uses both cpu and memory bandwidth) than malloc() because it must zero out memory it allocates at runtime
- Use calloc() only when you need the buffer to be zero filled prior to FIRST use

Introduction to Structs – An Aggregate Data Type

- **Structs** are a **collection (or aggregation) of values** grouped under a **single name**
 - Each **variable** in a struct is called a **member** (sometimes **field** is used)
 - Each **member** is identified with a **name**
 - Each **member** can be (and quite often are) **different types**, include other structs
 - Like a Java class, but no associated methods or constructors with a struct
- Structure definition **does not** define a variable instance, nor does it allocate memory:
 - It creates a **new variable type** uniquely identified by its **tagname**:
"struct tagname" includes the **keyword struct** and the **tagname** for this type

Easy to forget
semicolon!


```
struct tagname {  
    type1 member1;  
    ...  
    typeN memberN;  
};
```

```
struct vehicle {  
    char *state; 4  
    char *plate; 4  
    char *make; 4  
    int year; 4  
};
```

16 bytes

Struct Variable Definitions

```
struct vehicle {  
    char *state;  
    char *plate;  
    char *make;  
    int year;  
};  
struct vehicle name1;
```



- Variable definitions like any other data type:

```
struct vehicle name1, *pn, ar[3];
```

type: "struct vehicle"

single variable instance

pointer

array

Accessing members of a struct

- Like arrays, struct variables are aggregated contiguous objects in memory
- The `.` structure operator which *"selects"* the requested field or member

```
struct date { // defining struct type
    int month;
    int day; // members date struct
};
```

day	
month	

```
struct date bday; // struct instance
bday.month = 1;
bday.day = 24;

// shorter initializer syntax
struct date new_years_eve = {12, 31};
struct date final = {.day= 24, .month= 1};
```

struct date bday

day	24
month	1

Accessing members of a struct with pointers

```
struct date { // defining struct type
    int month;
    int day; // members date struct
};
```

day	
month	

- Now create a *pointer* to a struct

```
struct date *ptr = &bday;
```

- Two options to reference a member via a struct pointer (. is higher precedence than *):

- Use * and . operators: `(*ptr).month = 11;`

- Use -> operator for shorthand: `ptr->month = 11;`

*(* ptr). month*

Operator	Description	Associativity
() [] . -> ++ --	Parentheses or function call Brackets or array subscript Dot or Member selection operator Arrow operator Postfix increment/decrement	left to right
++ -- + - ! ~ (type) * & sizeof	Prefix increment/decrement Unary plus and minus not operator and bitwise complement type cast Indirection or dereference operator Address of operator Determine size in bytes	right to left

Accessing members of a struct

```
struct date {           // defining struct type
    int month;
    int day;            // members date struct
};
```

- You can create an array of structs and initialize them

```
struct date quarter[] =
    { {1,2}, {3,4}, {5,6}, {7,8}, {9,10} };
int cnt = sizeof(quarter)/sizeof(*quarter); // = 5
```

quarter[4]	day	10
	month	9
quarter[3]	day	8
	month	7
quarter[2]	day	6
	month	5
quarter[1]	day	4
	month	3
quarter[0]	day	2
	month	1


Accessing members of a struct

```
struct date quarter[3];  
struct date *ptr;
```

```
ptr = quarter + 1; // array name = address  
ptr->month = 2;  
ptr->day = 21; // or (*ptr).day = 21;
```

```
(ptr-1)->month = 1; // or (*(ptr-1)).month = 4;  
(ptr-1)->day = 7;
```

```
(++ptr)->month = 3;  
ptr->day = 5;
```



Operator	Description	Associativity
() [] . -> ++ --	Parentheses or function call Brackets or array subscript Dot or Member selection operator Arrow operator Postfix increment/decrement	left to right
++ -- + - ! ~ (type) * & sizeof	Prefix increment/decrement Unary plus and minus not operator and bitwise complement type cast Indirection or dereference operator Address of operator Determine size in bytes	right to left

Typedef usage with Struct – Another Style Conflict

- *Typedef* is a way to create an *alias* for another data type (not limited to just structs)

`typedef <data type> <alias>;`

- After typedef, the alias can be used interchangeably with the original data type

- e.g., `typedef unsigned long int size_t;`

- *Many claim typedefs* are easier to understand than tagged struct variables

⇒ *typedef with structs* are not allowed in the cse30 style guidelines (Linux kernel standards)

```
struct nm {  
    /* fields */  
};  
typedef struct nm item;  
  
item n1;  
struct nm n2;  
item *ptr;  
struct nm *ptr2;
```

```
typedef struct name2_s {  
    int a;  
    int b;  
} name2_s;  
  
name2_s var2;  
name2_s *ptr2;
```

```
typedef struct {  
    int a;  
    int b;  
} pair;  
  
pair var3;  
pair *ptr3;
```

Copying Structs

- You can assign the member value(s) of the whole struct from a struct of the same type
– *this copies the entire contents!*


- Individual fields can also be copied

```
struct date first = {1, 1};  
struct date final = {.day= 31, .month= 12};
```

```
struct date *pt1 = &first;  
struct date *pt2 = &final;
```

```
final.day = first.day; // both day are 1  
final = first;         // copies whole struct
```

```
pt2->month = 3;  
*pt1 = *pt2;           // copies whole struct  
pt2->day = 7;  
pt1->day = pt2->day;   // both days are now 7
```



Struct: Copy and Member Pointers

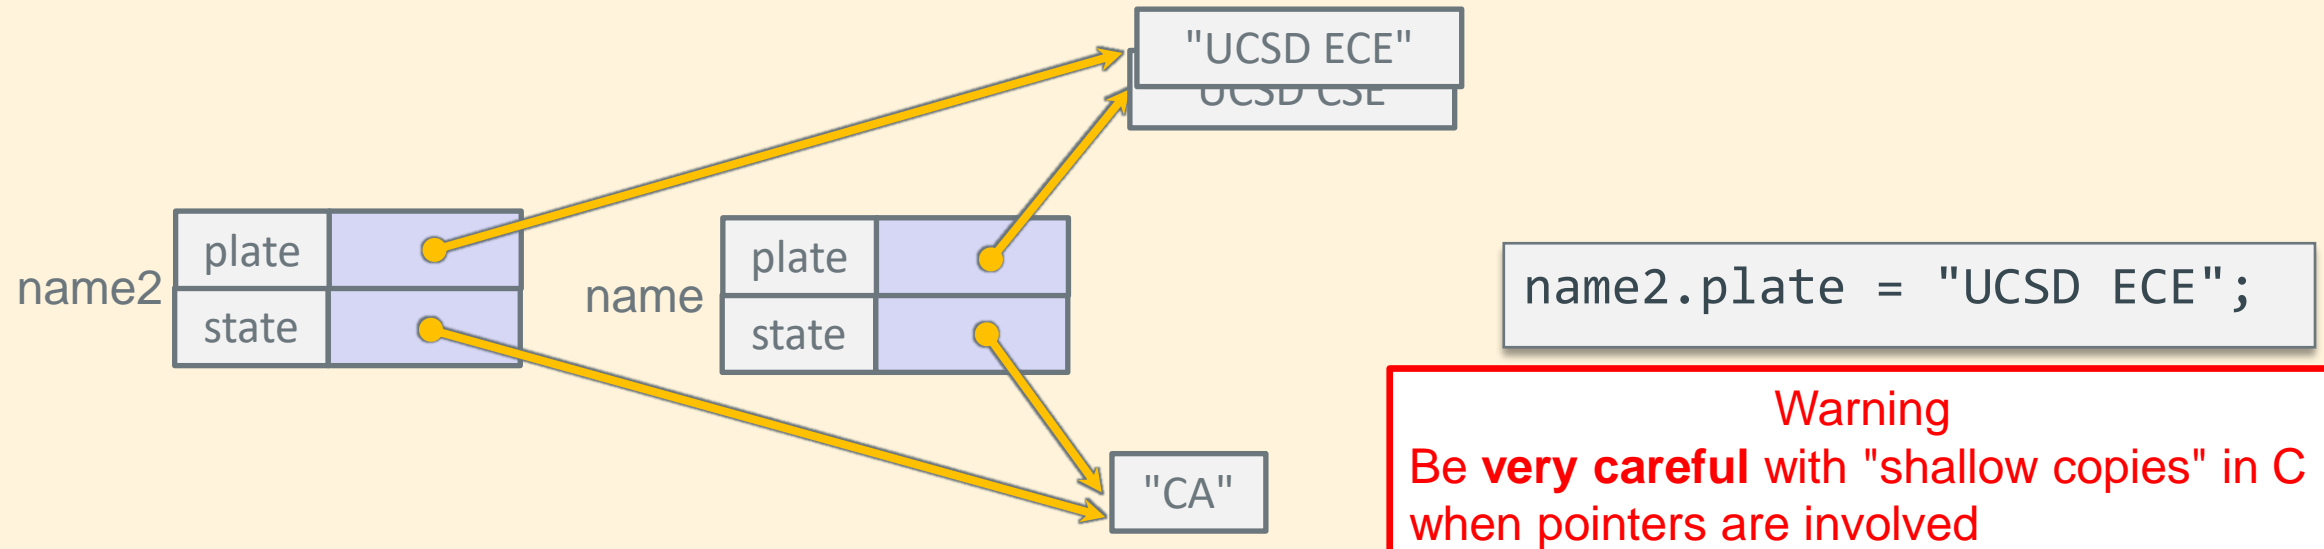
```
struct vehicle {  
    char *state;  
    char *plate;  
};
```

```
struct vehicle name = {"CA", "UCSD CSE"};  
struct vehicle name2;
```



- When you assign one struct to another it just copies the member fields!

```
name2 = name; // copies members Only
```




Memory Allocation Structs with Pointer Members

- **Safety first:** Allocate anything that is pointed at by a struct member independently (they are not part of the struct, only the pointers are)

```
struct vehicle {  
    char *state;  
    char *plate;  
    int year;  
};  
struct vehicle name1;  
pn = &name1;
```

```
name1.state = strdup("CA");  
pn->plate = strdup("xyz");  
pn->year = 2021;
```



22

X

X

X

- X

X



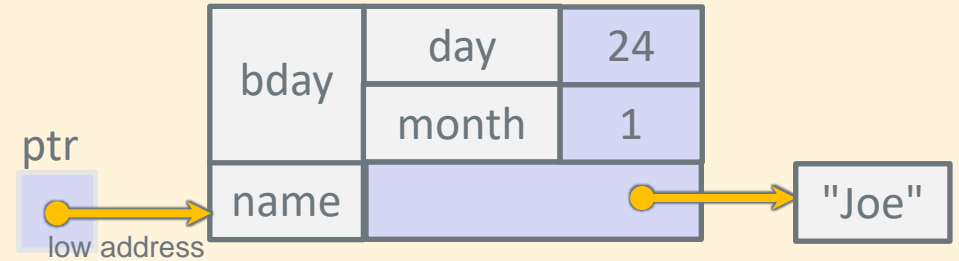
X

Nested Structs

- Structs like any other variable can be a member of a struct, this is called a **nested struct**

```
struct date {  
    int month;  
    int day;  
};
```

```
struct person {  
    char *name;  
    struct date bday;  
};
```

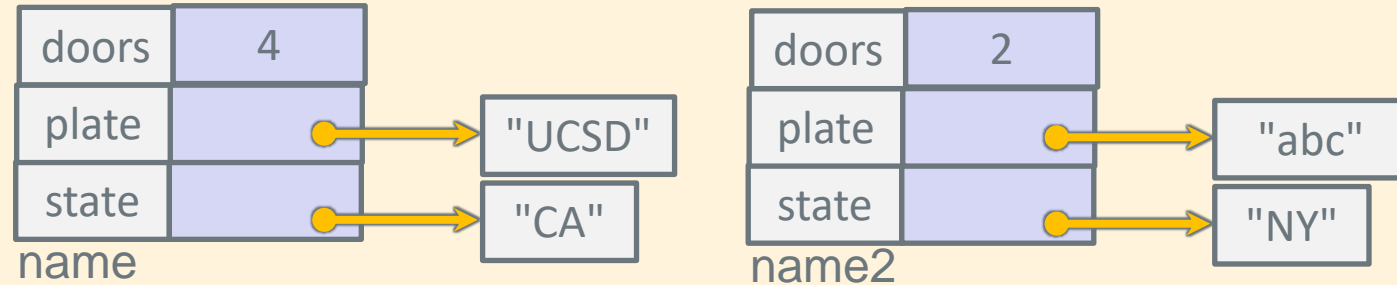


```
struct person first;  
struct person *ptr;  
ptr = &first;  
  
first.name = "Sam"; // immutable string  
first.name = (char []) {"Joe"}; // mutable string, lost address to Sam  
  
first.bday.month = 1;  
first.bday.day = 24;  
  
// below is the same as above  
ptr->bday.month = 1;  
ptr->bday.day = 24;
```

Comparing Two Structs

- You cannot compare entire structs, you must compare them one member at a time

```
struct vehicle {  
    char *state;  
    char *plate;  
    int doors;  
};
```



```
struct vehicle name = {"CA", "UCSD", 4};  
struct vehicle name2 = { (char []) {"NY"}, (char []) {"abc"}, 2};
```

```
if ((strcmp(name.state, name2.state) == 0) &&  
    (strcmp(name.plate, name2.plate) == 0) &&  
    (name.doors == name2.doors)) {  
    printf("Same\n");  
} else {  
    printf("Different\n");  
}
```


Struct: Arrays and Dynamic Allocation


- Allocate individual structs and arrays of structs using malloc()
 - Remember `.` is higher precedence than `*`:

```
#define HOLIDAY 5
struct date *pt1 = malloc(sizeof(*pt1));
struct date *pt2 = malloc(sizeof(*pt2) * HOLIDAY);
```

```
(*ptr1).month = 2;
(*ptr1).day = 21;

pt2->month = 12;
pt2->day = 25;
(pt2+1)->day = 22;    //or (*(pt2+1)).month
```

```
free(pt1);
pt1 = NULL;
free(pt2);
pt2 = NULL;
```




Struct As A Parameter to Functions

- **WARNING:** When you pass a struct, **you pass a copy** of all the struct members
 - This is a shallow copy (shallow copy – so if you have members that are pointers watch out)
- More often code will pass the pointer to a struct to **avoid the copy costs**
 - Be careful and not modify what the pointer points to (unless it is an output parameter)
- Tradeoffs:
 - Passing a pointer is cheaper and takes less space unless struct is small
 - **Member access cost:** indirect accesses through pointers to a struct member **may** be a bit more expensive and **might be harder for compiler to optimize**
- For small structs like a struct date passing a copy is fine
- **For** large structs always use pointers (arrays of struct, being an array is always a pointer)
 - For me, I always use pointers regardless of size, but that is just maybe a decades old habit...


Struct as a Parameter to Functions – Be Careful it is not like arrays

```
void change1(struct vehicle car)
{
    car.date = 2022; // oops!
    *(car.state) = "P";
}
...
change1(name);
```

Changes two different structs



```
void change2(struct vehicle *car)
{
    car->date = 2022;
    *(car->state) = "P";
}
...
change2(name);
```

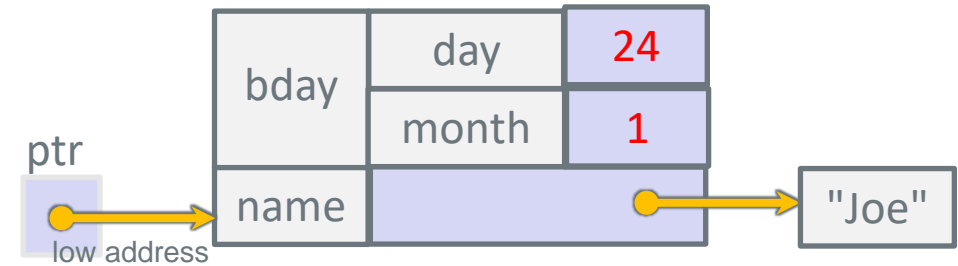


Changes the caller's version

Struct as an Output Parameter: Deep Copy Example

```
struct date {  
    int month;  
    int day;  
};
```

```
struct person {  
    char *name;  
    struct date bday;  
};
```




```
int fill(struct person *ptr, char *name, int month, int day)  
{  
    ptr->bday.month = month;  
    ptr->bday.day = day;  
    if ((ptr->name = strdup(name)) == NULL)  
        return -1;  
    return 0;  
}  
  
...-----calling function -----  
    struct person first;  
    if (fill(&first, "Joe", 1, 24) == 0)  
        printf("%s %d %d\n", first.name, first.bday.month, first.bday.day);  
    ...
```

Review: Singly Linked List - 1



- Is a **linear collection of nodes** whose order is not specified by their relative location in memory, like an array
- Each node consists of a **payload** and a **pointer** to the next node in the list
 - The **pointer in the last node** in the list is **NULL** (or 0)
 - The **head pointer points at the first node** in the list (the head is not part of the list)
- Nodes are **easy to insert and delete** from any position **without having to re-organize the entire data structure**
- Advantages of a linked list:
 - **Length can easily be changed** (expand and contract) at execution time
 - **Length does not need to be known in advance** (like at compile time)
 - List can **continue to expand** while there is memory available

Review: Singly Linked List - 2




- Memory for each node is typically **allocated dynamically at execution time** (*i.e.*, using *heap memory* – *malloc()* *etc.*) when a new node is added to the list
- Memory for each node may be freed at execution time, using *free()* **when a node is removed from the list**
- Unlike arrays, linked **list nodes are usually not arranged** (located) sequentially in adjacent memory locations
- **No fast and convenient way** to "jump" to any specific node.
- Usually the list must be **traversed (walked)** from the **head** to locate if a **specific payload** is stored in any node
- Obviously, the cost in **traversing a linked list is $O(n)$**

Linked List Using Self-Referential Structs

- A **self-referential struct** is a struct that has one or more **members** that are **pointers** to a **struct variable of the same type**

- Self-referential member
 - points to same type – itself


```
struct node {  
    int date;  
    struct node *next;  
};
```



- Example:

```
struct node2 {  
    int month;  
    int day;  
    struct node2 *next;  
};  
struct node2 x;
```

```
struct node2 *head  
head = &x;
```



```
x.month = 1;  
x.day = 31;  
x.next = NULL;
```

Creating a Node & Inserting it at the Front of the List



```
// create node; insert at front when passed head
struct node *
creatNode(int year, char *name, struct node *link)
{
    struct node *ptr = malloc(sizeof(*ptr));
    if (ptr != NULL) {
        if ((ptr->name = strdup(name)) == NULL) {
            free(ptr);
            return NULL;
        }
        ptr->year = year;
        ptr->next = link;
    }
    return ptr;
}
```

Must duplicate
the string
because of
buffer reuse

```
struct node {
    int year;
    char *name;
    struct node *next;
};
```

```
// calling function body
struct node *head = NULL; // insert at front
struct node *ptr;
char buf[BUFSZ];

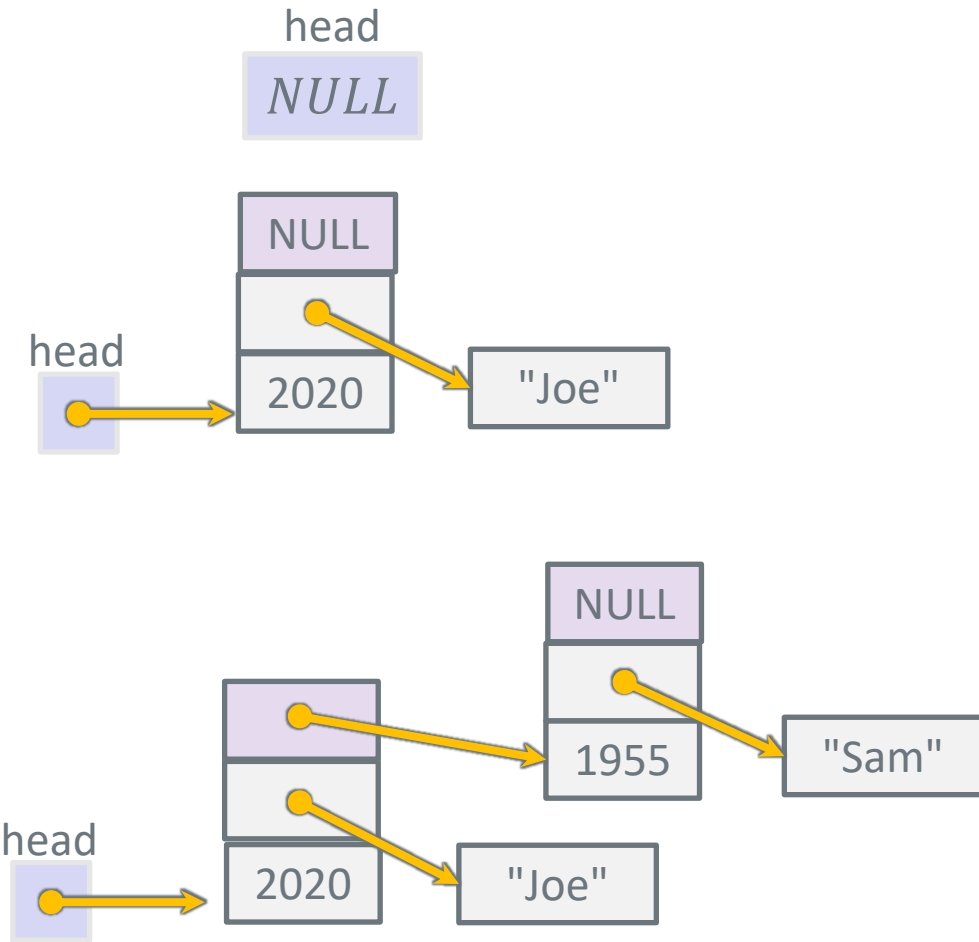
if (fgets(buf, BUFSZ, stdin) != NULL) { // reads joe
    if ((ptr = creatNode(2020, buf, head)) != NULL) {
        head = ptr; // error handling not shown
    }
}
if (fgets(buf, BUFSZ, stdin) != NULL) { // reads sam
    if ((ptr = creatNode(1955, buf, head)) != NULL) {
        head = ptr; // error handling not shown
    }
}
```



Creating a Node & Inserting it at the **End** of the List

```
// create a node and insert at the end of the list
```

```
struct node *  
insertEnd(int year, char *name, struct node *head)  
{  
    struct node *ptr = head;  
    struct node *prev = NULL; // base case  
    struct node *new;  
  
    if ((new = creatNode(year, name, NULL)) == NULL)  
        return NULL;  
  
    while (ptr != NULL) {  
        prev = ptr;  
        ptr = ptr->next;  
    }  
    if (prev == NULL)  
        return new;  
    prev->next = new;  
    return head;  
}
```




```
struct node *head = NULL; // insert at end  
struct node *ptr;  
if ((ptr = insertEnd(2020, "Joe", head)) != NULL)  
    head = ptr;  
if ((ptr = insertEnd(1955, "Sam", head)) != NULL)  
    head = ptr;
```

"Dumping" the Linked List

"walk the list from head to tail"


```
struct node {  
    int year;  
    char *name;  
    struct node *next;  
};
```



```
struct node *head;  
struct node *ptr;  
...  
printf("\nDumping All Data\n");  
ptr = head;  
while (ptr != NULL) {  
    printf("year: %d name: %s\n", ptr->year, ptr->name);  
    ptr = ptr->next;  
}
```

Dumping All Data
year: 1955 name: Sam
year: 2020 name: Joe

Finding A Node Containing a Specific Payload Value



```
struct node {
    int year;
    char *name;
    struct node *next;
};
```

```
struct node *findNode(char *name, struct node *ptr)
{
    while (ptr != NULL) {
        if (strcmp(name, ptr->name) == 0)
            break;
        ptr = ptr->next;
    }
    return ptr;
}
```

Returns pointer if found
NULL otherwise

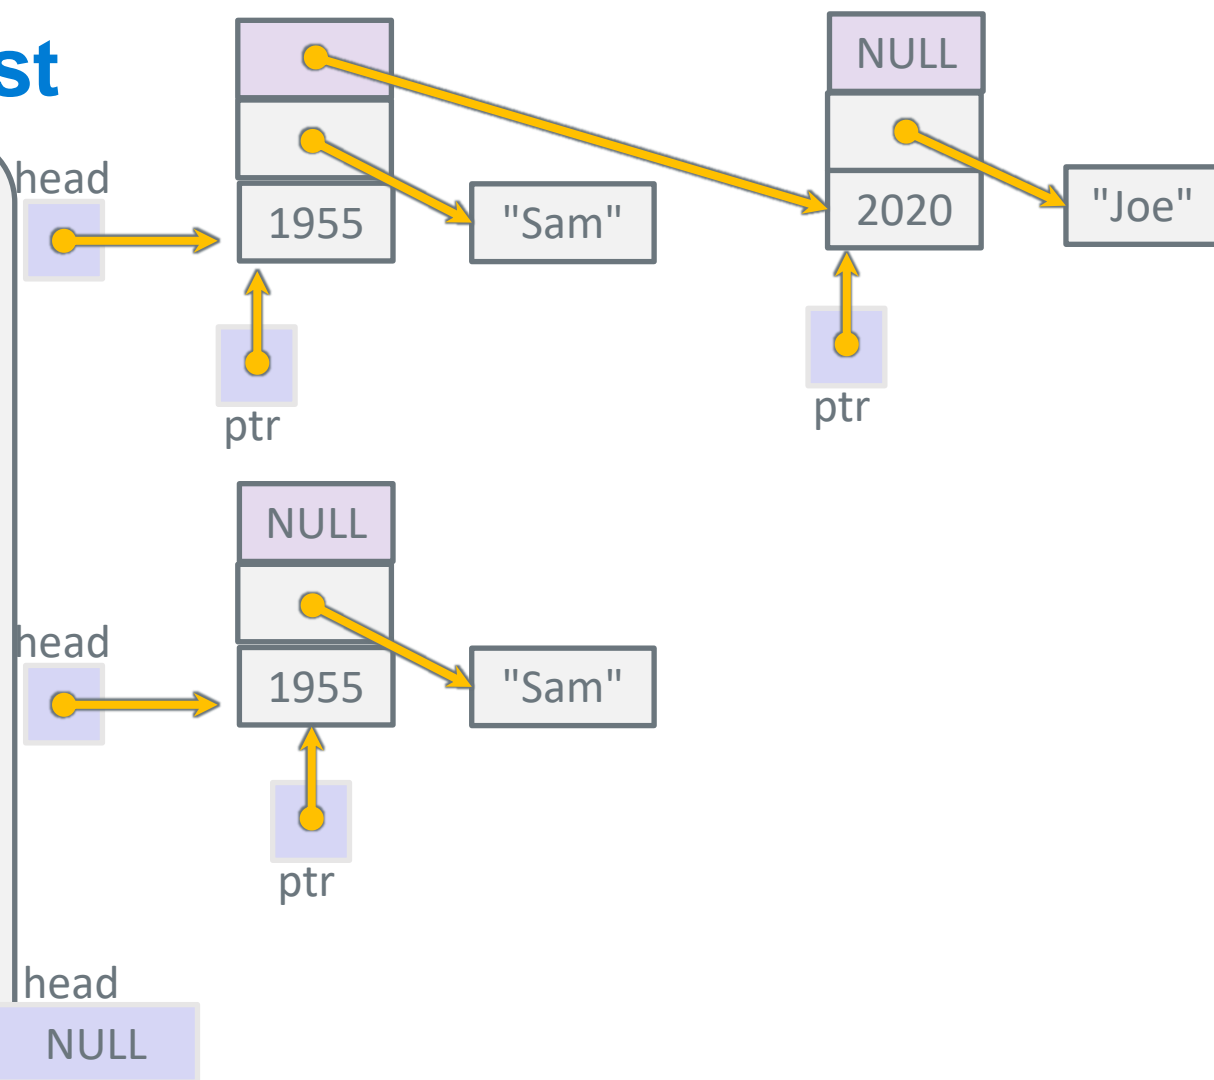
```
struct node *found;
```

```
if ((found = findNode("Joe", head)) != NULL)
    printf("year: %d name: %s\n", found->year, found->name);
```

Deleting a Node in a Linked List


```
// returns head pointer; may have changed...
struct node *deleteNode(int name, struct node *head)
{
    struct node *ptr = head;
    struct node *prev = NULL;

    while (ptr != NULL) {
        if (strcmp(name, ptr->name) == 0)
            break;
        prev = ptr;
        ptr = ptr->next;
    }
    if (ptr == NULL) // not found return head
        return head;
    if (ptr == head) // remove first node new head
        head = ptr->next;
    else
        prev->next = ptr->next; // remove not head
    free(ptr->name); // free strdup() space
    free(ptr);
    return head;
}
```





```
struct node *head = NULL;
head = deleteNode("Joe", head);
head = deleteNode("Sam", head);
```

Improving On Linked List Performance




- When linked lists get long, the cost of finding an entry continues to increase $O(n)$
- How to improve search time?
- Break the single linked list into multiple **shorter length** linked lists
 - Shorter lists are faster to search
- **Requires a function** that takes a **lookup key** and selects just one of the shortened lists

How do you determine on which linked list an entry is stored?




same as →




Hashing

- Hash table is an array of **pointers** to the head of different linked lists (called **hash chains**)
- **Each data item** must have a **unique key** that **identifies it** (e.g., **auto license plate**)
 - $h(k)$ is the **hash value** of key k to **encode the key k into an integer**
- Use the Hash value to map to **one entry** in the hash table $T[]$ of size TABLESIZE
 - $\text{Index} = h(k) \% \text{TABLESIZE}$ (mod operator $\%$ maps a **keys hash value** to **table index**)
- **Keys** that hash to the same array index (**collide**) are **put on a linked list**
- After hashing a **key**, you then traverse the selected linked list to find the entry



Hash Table With Collision Chaining (multiple linked lists)



- Make $TABLESIZE$ **prime** as keys are typically not randomly distributed, and have a *pattern*
 - Mostly even, mostly multiples of 10, etc.
 - In general: mostly multiples of some k
- If k is a **factor** of $TABLESIZE$, then only $(TABLESIZE/k)$ slots will ever be used!

1. Calculate index $i = \text{hash}(\text{key}) \% TABLESIZE$
 2. Go to array element i , i.e., $T+i$ that contains the head pointer for collision chain
 3. Walk the linked list for element, add element, remove element, etc. from the linked list
 4. New items added to the hash table are typically added at the front or at the end of the *collision chain* linked list (*when multiple keys hash to same index .. they collide*)
- Hash arrays need an index number to select a chain, so if we have a string, we must first convert to a number


Simple 32-bit String Hash Function in C (djb2)

```
uint32_t hash(char *str)
{
    uint32_t hash = 0U;
    uint32_t c;

    while ((c = (unsigned char)*str++) != '\0')
        hash = c + (hash << 6) + (hash << 16) - hash;

    return hash;
}
```

Signed Data types	Unsigned Data types	Exact Size
int32_t	uint32_t	32 bits (4 bytes)



- Many different algorithms for string hash function (Dan Berman's djb2 above)
 - << is the **left** bit shift operator (later in course)
- **Fast to compute**, has a reasonable key distribution for **short 8-bit ASCII strings** into **32-bit unsigned ints**

Allocating the Hash Table (collision chain head pointers)

Good use for calloc()

```
#define TBSZ 3
int main(void)
{
    struct node *ptr;
    struct node **tab;
    uint32_t index;

    if ((tab = calloc(TBSZ, sizeof(*tab))) == NULL) {
        fprintf(stderr, "Cannot allocate hash table\n");
        return EXIT_FAILURE;
    }
    // continued on next slide
```



TABLESIZE = 3

Inserting Nodes into the Hash Table (at the end)

```
#define TBSZ 3
unit32_t index;

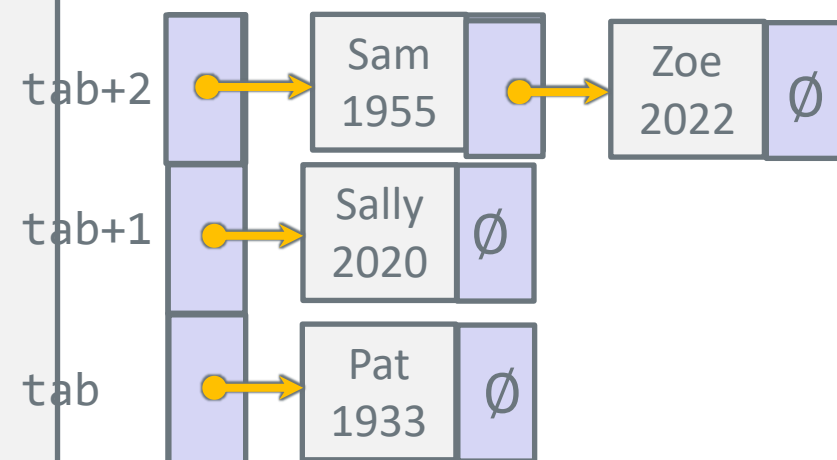
index = hash("Pat") % TBSZ;
if ((ptr = insertEnd(1933, "Pat", *(tab + index))) != NULL)
    *(tab + index) = ptr;

index = hash("Sam") % TBSZ;
if ((ptr = insertEnd(1955, "Sam", *(tab + index))) != NULL)
    *(tab + index) = ptr;

index = hash("Sally") % TBSZ;
if ((ptr = insertEnd(2020, "Sally", *(tab + index))) != NULL)
    *(tab + index) = ptr;

index = hash("Zoe") % TBSZ;
if ((ptr = insertEnd(2022, "Zoe", *(tab + index))) != NULL)
    *(tab + index) = ptr;
```


```
struct node {
    int year;
    char *name;
    struct node *next;
};
```



Notice

Substitute **createNode()** for **insertEnd()** to insert nodes at the **front** of the collision chain **instead** of at the **end** of the collision chain


"Dumping" the Hash Table (traversing all Nodes)



Dumping All Data
chain: 0
Year: 1933 Name: Pat
chain: 1
Year: 2020 Name: Sally
chain: 2
Year: 1955 Name: Sam
Year: 2022 Name: Zoe

```
printf("\nDumping All Data\n");  
for (index = 0U; index < TBSZ; index++) {  
    ptr = *(tab + index);  
    printf("chain: %d\n", index);  
  
    while (ptr != NULL) {  
        printf("Year: %d Name: %s\n", ptr->year, ptr->name);  
        ptr = ptr->next;  
    }  
}
```

Finding a Node with a Specific Payload Value



```
// same routine as shown in a previous slide
struct node *findNode(char *name, struct node *ptr)
{
    while (ptr != NULL) {
        if (strcmp(name, ptr->name) == 0)
            break;
        ptr = ptr->next;
    }
    return ptr;
}
```

```
index = hash("Zoe") % TBSZ;
if ((ptr = findNode("Zoe", *(tab + index))) != NULL)
    printf("Found Year: %d name: %s\n", ptr->year, ptr->name);
else
    printf("Not Found Zoe\n");
```

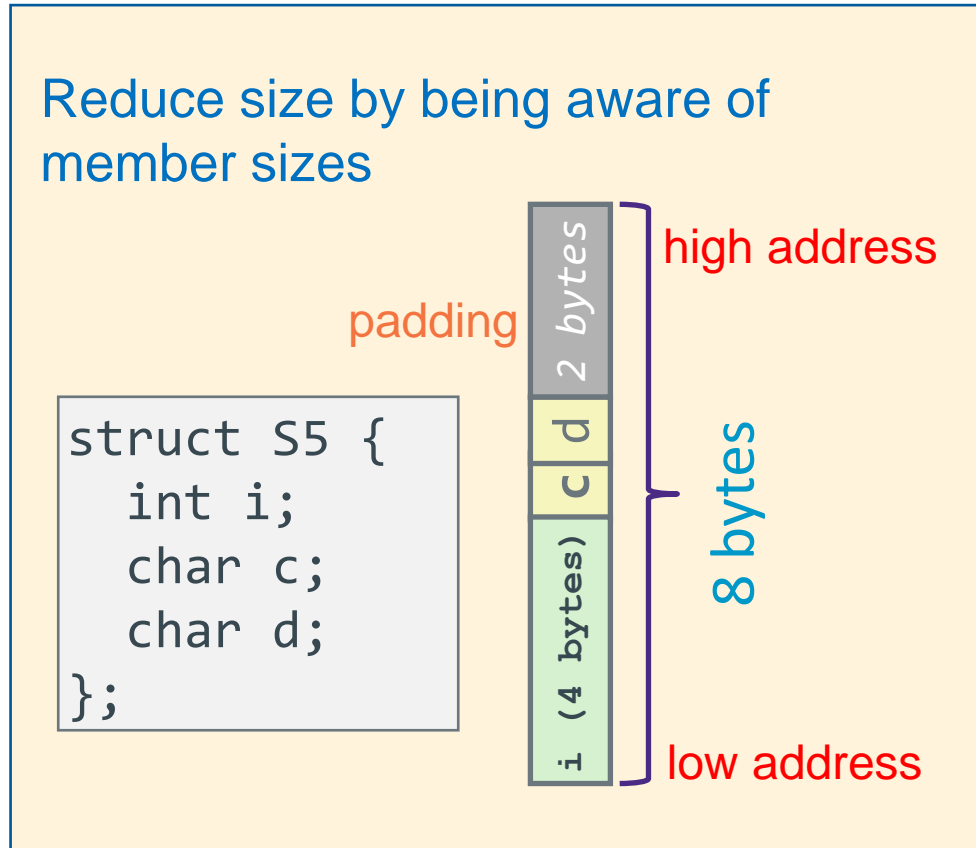
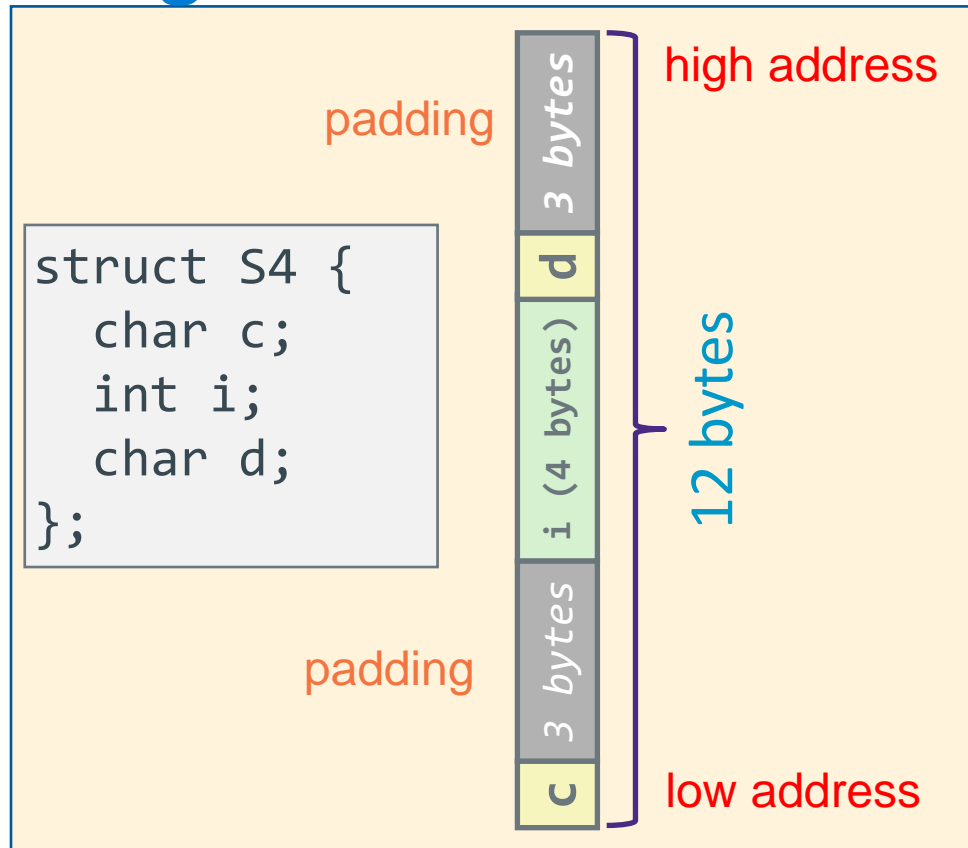
Sizing Struct Members

- struct size depends on the order of the fields listed in the struct

```
struct S4 {  
    char c;    // byte aligned  
    int i;     // 4 byte aligned  
    char d;    // byte aligned  
};
```

- Structs uses contiguously-allocated region of memory,
- compilers **are required to follow member order**, and HW alignment requirements
 1. not allowed to re-arrange member order in memory
 2. struct starting address: aligned to the requirements of largest member
 3. Add memory space between members (pad or unused space), so the next member starts at the required memory alignment
 4. Structs may add padding so total size is always a whole multiple of the size of the largest member (for struct arrays)

Re-Sizing Struct Members



- re-order the fields to decrease space wasted by member alignment padding
- Remember by C specifications, the compiler will not do this for you...
- To get the byte offset (from the start) of any member of a struct

```
#include <stddef.h>  
size_t cnt = offsetof(struct_name, member_name);
```